



UNIVERSITY OF TARTU

INSTITUTE OF COMPUTER SCIENCE



Basics of Cloud Computing – Lecture 3

Scaling Applications on the Cloud

Satish Srirama



Mobile & Cloud Lab

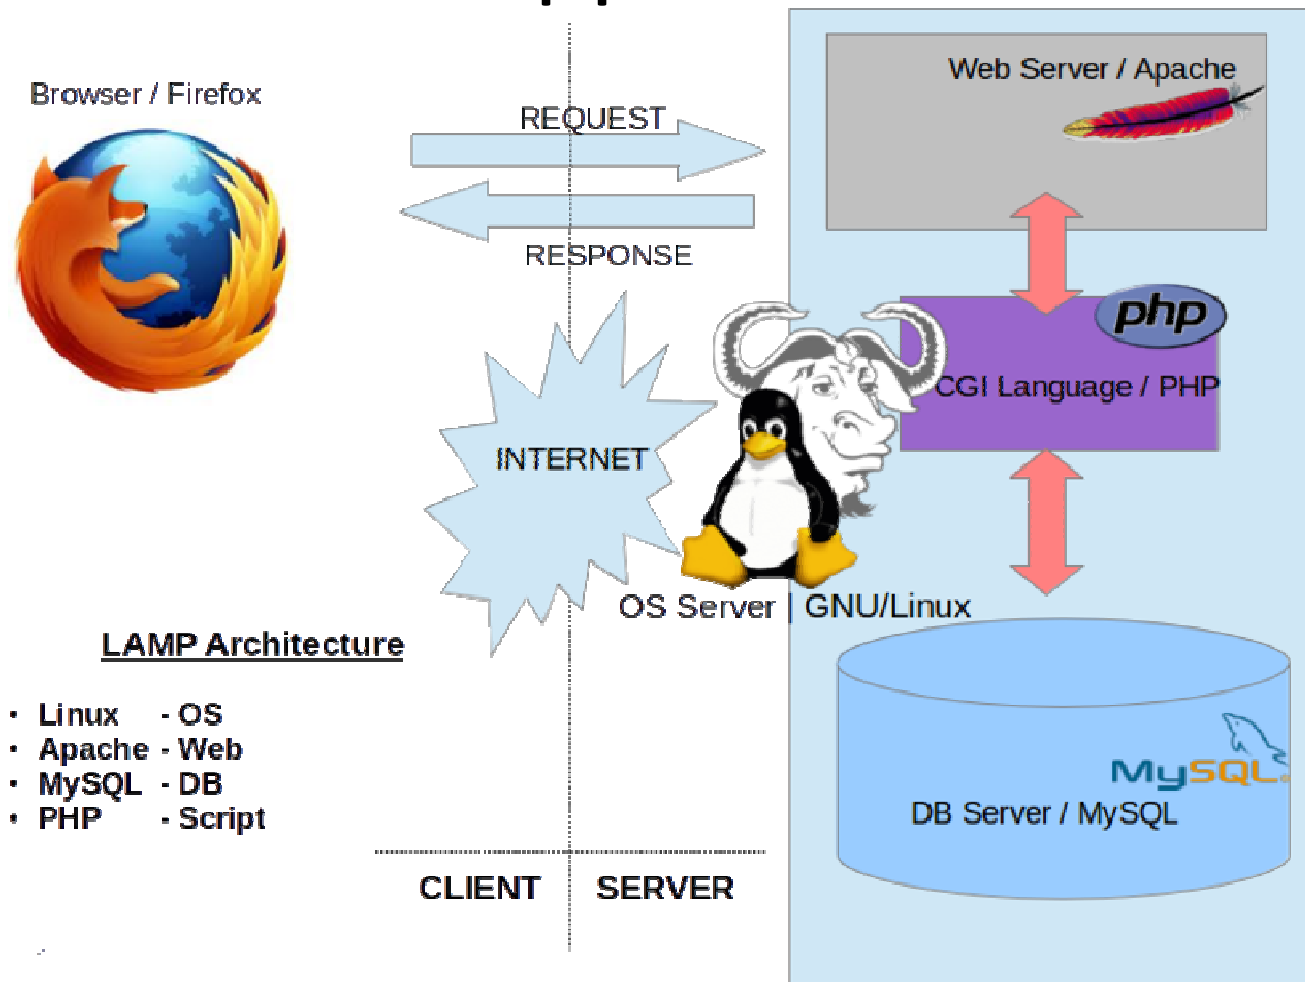
Outline

- Scaling Information Systems
- Scaling Enterprise Applications in the Cloud
- Auto Scaling
 - Amazon Auto Scale and Elastic Load Balancing
 - Open source options for developing Auto Scale solutions

Scaling Information Systems


- Fault tolerance, high availability & scalability are essential prerequisites for any enterprise application deployment
- Scalability
 - Generally nodes in information systems support specific load
 - When load increases beyond certain level, systems should be scaled up
 - Similarly when load decreases they should be scaled down

Typical Web-based Enterprise Application



Source: http://en.wikipedia.org/wiki/File:LAMPP_Architecture.png

Typical Load on an Application Server




[ClarkNet traces]

Scaling Information Systems - continued

- Two basic models of scaling
 - Vertical scaling
 - Also known as Scale-up
 - Horizontal scaling
 - Aka Scale-out


Vertical Scaling

- Achieving better performance by replacing an existing node with a much powerful machine
- Risk of losing currently running jobs
 - Can frustrate customers as the service is temporarily down



Horizontal Scaling

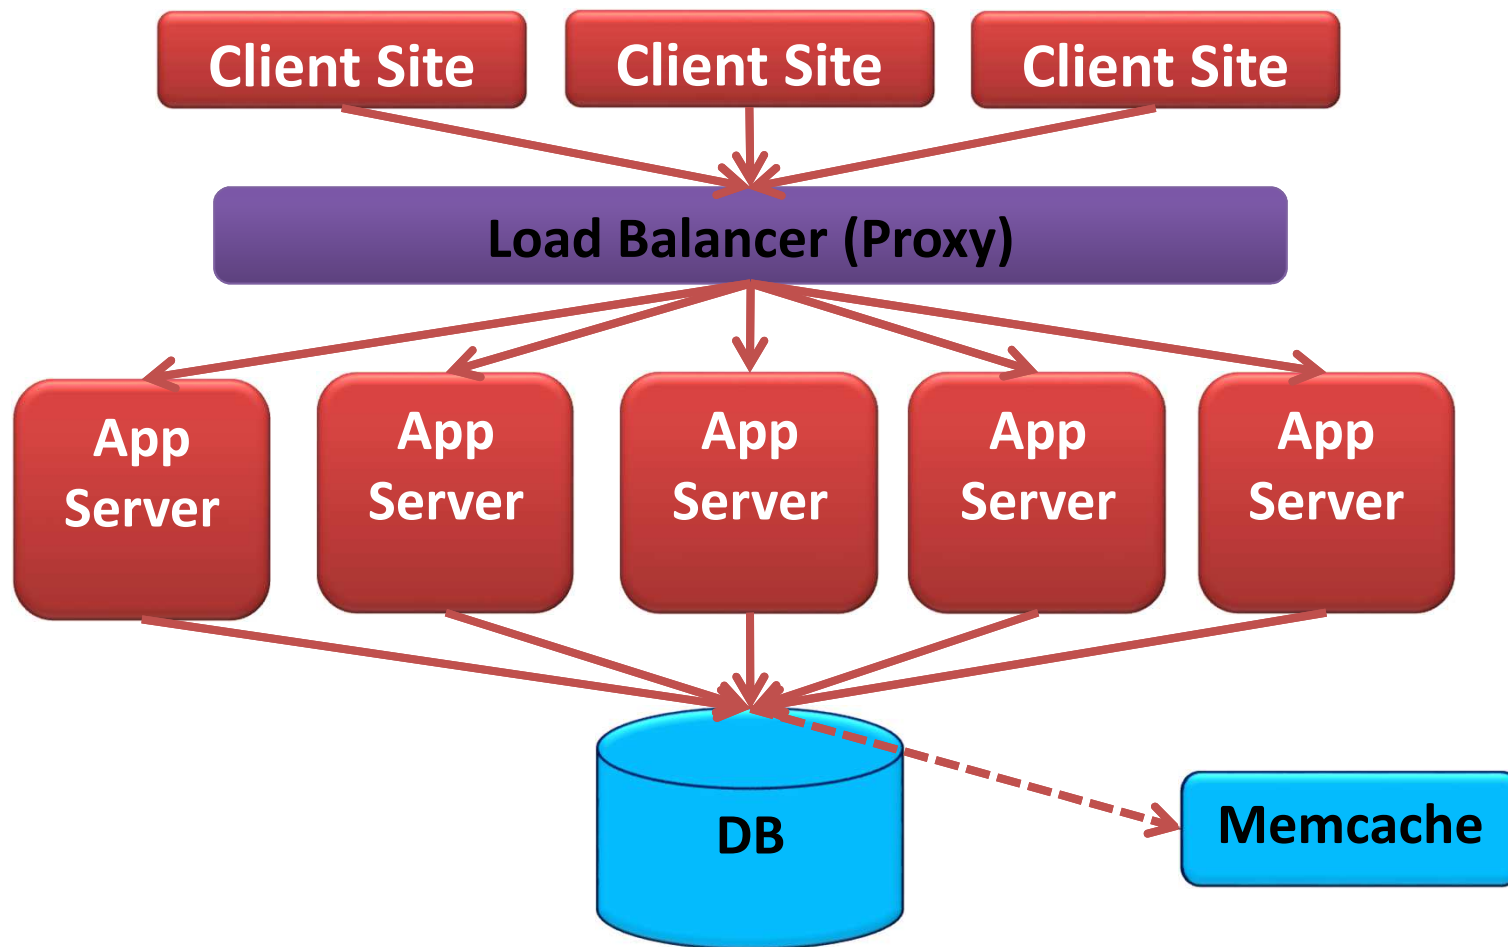
- Achieving better performance by adding more nodes to the system
- New servers are introduced to the system to run along with the existing servers



Server allocation policies for different loads

[Vasar et al, Nordcloud 2012]

Scaling Enterprise Applications in the Cloud



Load Balancer

- Load balancing has been a key mechanism in making efficient web server farms
- Load balancer automatically distributes incoming application traffic across multiple servers
- Hides the complexity for content providers
- $1+1 = 2$
 - Allows server farms work as a single virtual powerful machine
- $1+1 > 2$
 - Beyond load distribution, improves response time

Introduction- Types of Load Balancers

- Network-Based load balancing
 - Provided by IP routers and DNS (domain name servers) that service a pool of host machines
 - e.g. when a client resolves a hostname, the DNS can assign a different IP address to each request dynamically based on current load conditions
- Network-Layer based load balancing
 - Balances the traffic based on the source IP address and/or port of the incoming IP packet
 - Does not take into account the contents of the packet, so is not very flexible
- Transport-Layer based load balancing
 - The load balancer may choose to route the entire connection to a particular server
 - Useful if the connections are short-lived and are established frequently
- Application-Layer/Middleware based load balancing
 - Load balancing is performed in the application-layer, often on a per-session or per-request basis

Introduction- Classes of Load Balancers

- Non-adaptive load balancer
 - A load balancer can use non-adaptive policies, such as simple round-robin algorithm, hash-based or randomization algorithm
- Adaptive load balancer
 - A load balancer can use adaptive policies that utilize run-time information, such as amount of CPU load on the node
- Load Balancers and Load Distributors are not the same thing
 - Strictly speaking non-adaptive load balancers are load distributors

Load Balancing Algorithms


- Random
 - Randomly distributes load across the available servers
 - Picks one via random number generation and sending the current connection to it
- Round Robin
 - Round Robin passes each new connection request to the next server in line
 - Eventually distributes connections evenly across the array of machines being load balanced
- Least connection (Join-Shortest-Queue)
 - The system passes a new connection to the server that has the least number of current connections
- etc.
<https://devcentral.f5.com/articles/intro-to-load-balancing-for-developers-ndash-the-algorithms#.UwSBpc4a5SM>

Examples of Load Balancers

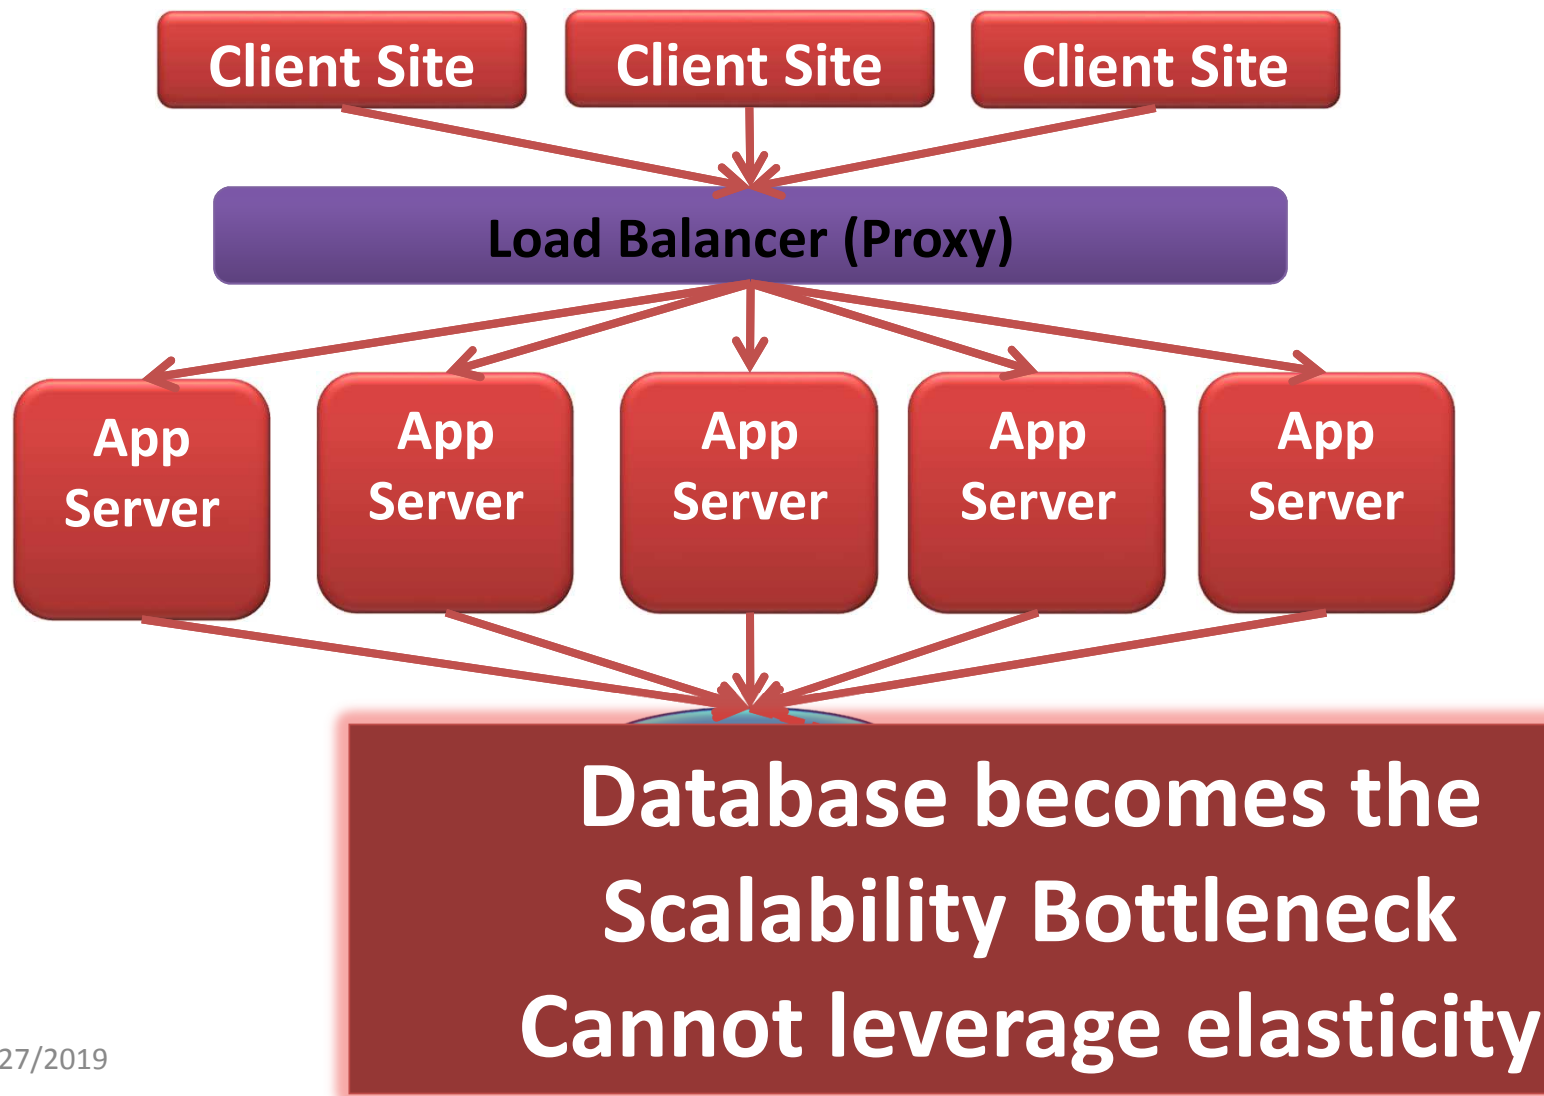
- Nginx - <http://nginx.org/>
 - HAProxy - <http://haproxy.1wt.eu/>
 - Pen - <http://siag.nu/pen/>
- etc.

Testing the System by Simulating Load

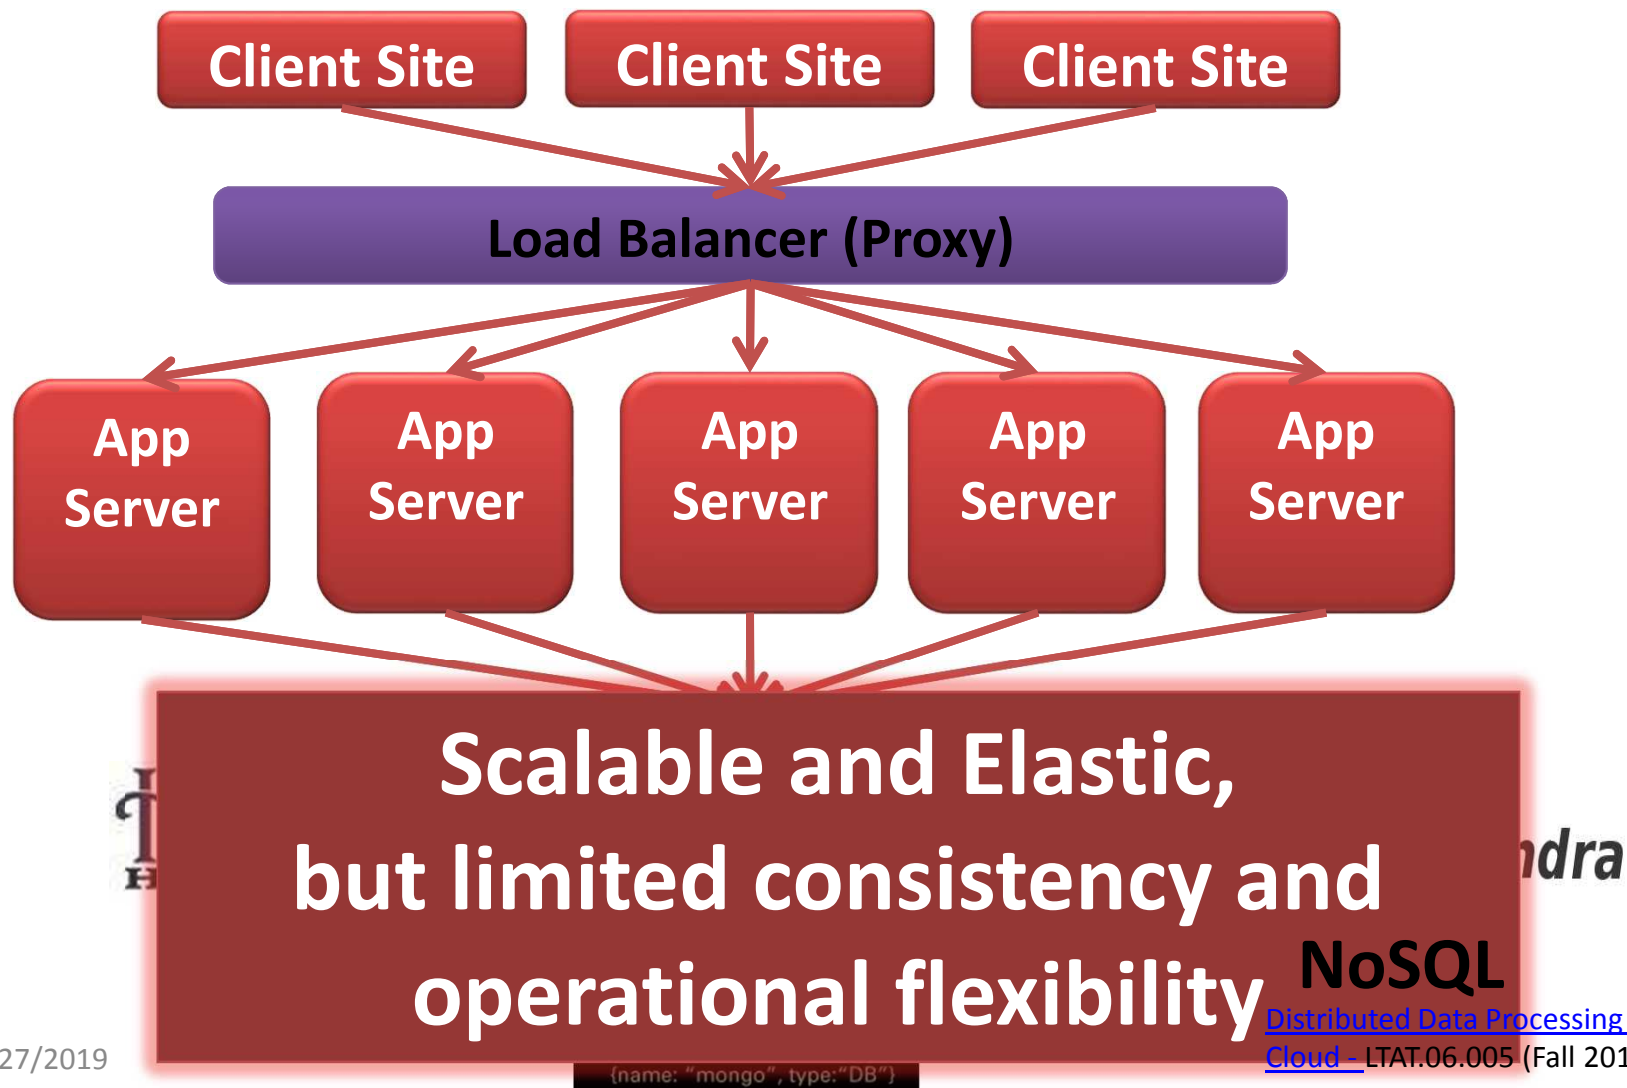
- Benchmarking tools
 - Tsung, JMeter, etc
- Simulating concurrency is also possible
- Multiple protocols
 - HTTP, XMPP, etc.
 - SSL support



Scaling in the Cloud - bottleneck



Scaling in the Cloud - bottleneck

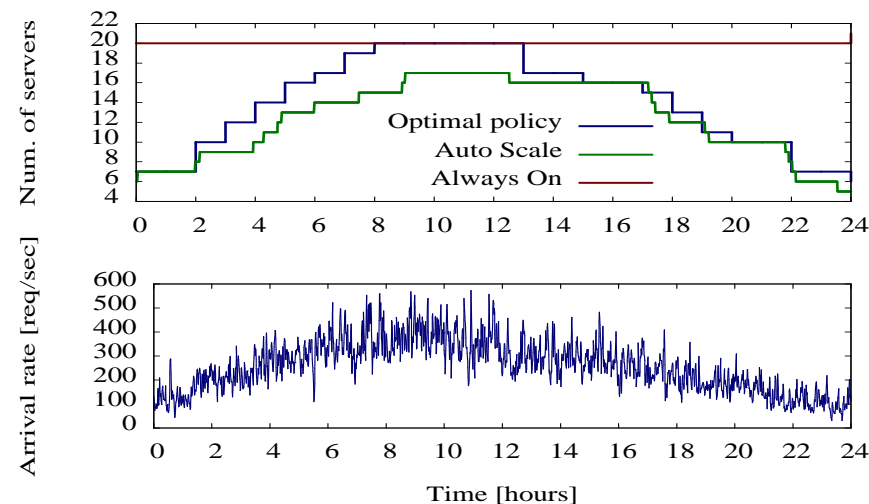


Horizontal Scaling – Further examples

- MapReduce & Hadoop
 - We will look as part of Lectures 4 and 5

AutoScale

- AutoScale allows systems to dynamically react to a set of defined metrics and to scale resources accordingly
- Providing:
 - High availability
 - Cost saving
 - Energy saving







Server allocation policies for different loads

Typical Usecases

- Applications that see elasticity in their demand
- Launching a new website with unknown visitor numbers
- Viral marketing campaigns
- A scientific application might also have to scale out
 - Using 50 machines for 1 hour rather than 1 machine for 50 hours

Types of Traffic Patterns


- ON & OFF
 - Analytics!
 - Banks/Tax Agencies!
 - Test environments
- FAST GROWTH
 - Events!
 - Business Growth!
- VARIABLE
 - News & Media!
 - Event Registrations!
 - Rapid fire sales
- CONSISTENT
 - HR Application!
 - Accounting/Finance



<http://www.slideshare.net/lynxmanuk/autoscaling-best-practices>

Auto-Scaling enterprise applications on the cloud

- Enterprise applications are mostly based on SOA and componentized models
- Auto-Scaling
 - Scaling policy -> When to Scale
 - Resource provisioning policy -> How to scale
- Threshold-based scaling policies are very popular due to their simplicity
 - Observe metrics such as CPU usage, disk I/O, network traffic etc.
 - E.g. Amazon AutoScale, RightScale etc.
 - However, configuring them optimally is not easy



SOA - Service-oriented architecture

AutoScaling on the cloud

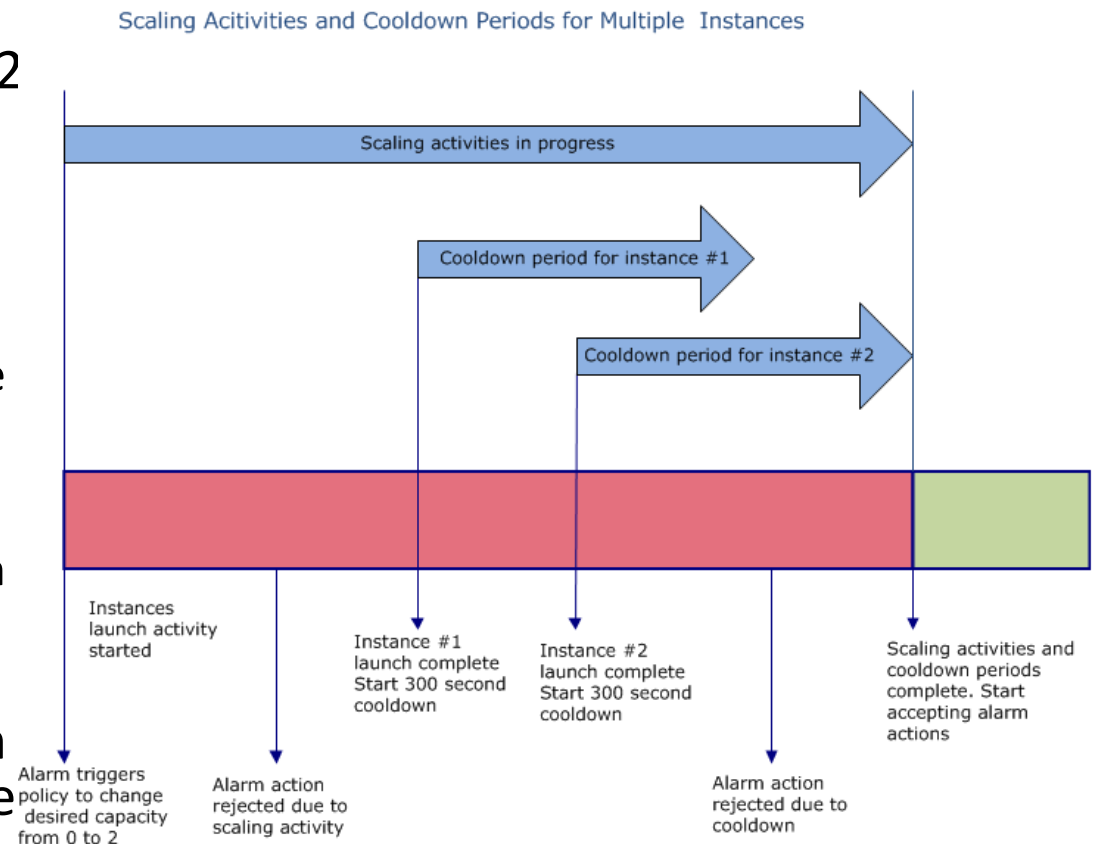
- Amazon Autoscale & Elastic Load Balance
- Vendor neutral autoscaling on cloud
 - Static Load Balancer + Resources estimation on the fly (e.g optimal heuristics) [Vasar et al, Nordiccloud 2012]
 - Static Load Balancer + Optimal resource provisioning [Srirama and Ostovar, CloudCom 2014]

Amazon Auto Scaling

- Amazon Auto Scaling allows you to scale your compute resources dynamically and predictably (*scaling plan*):
 - Dynamically based on conditions specified by you
 - E.g. increased CPU utilization of your Amazon EC2 instance
 - CPU utilization of all servers on average is >75% in last 5 min, add 2 servers and average < 35% remove 1 server
 - Predictably according to a schedule defined by you
 - E.g. every Friday at 13:00:00.
- EC2 instances are categorized into *Auto Scaling groups* for the purposes of instance scaling and management
- You create Auto Scaling groups by defining the minimum & maximum no of instances
- A launch configuration template is used by the Auto Scaling group to launch Amazon EC2 instances

Amazon Auto Scaling - continued

- Auto Scaling
 - Monitor the load on EC2 instances using CloudWatch
 - Define Conditions and raise alarms
 - E.g. Average CPU usage of the Amazon EC2 instances, or incoming network traffic from many different Amazon EC2 instances
 - Spawn new instances when there is too much load or remove instance when not enough load



Amazon CloudWatch

- Monitor AWS resources automatically
 - Monitoring for Amazon EC2 instances: seven pre-selected metrics at five-minute frequency
 - Amazon EBS volumes: eight pre-selected metrics at five-minute frequency
 - Elastic Load Balancers: four pre-selected metrics at one-minute frequency
 - Amazon RDS DB instances: thirteen pre-selected metrics at one-minute frequency
 - Amazon SQS queues: seven pre-selected metrics at five-minute frequency
 - Amazon SNS topics: four pre-selected metrics at five-minute frequency
- Custom Metrics generation and monitoring
- Set alarms on any of the metrics to receive notifications or take other automated actions
- Use Auto Scaling to add or remove EC2 instances dynamically based on CloudWatch metrics

Elastic Load Balance

- Elastic Load Balance
 - Automatically distributes incoming application traffic across multiple EC2 instances
 - Detects EC2 instance health and diverts traffic from bad ones
 - Support different protocols
 - HTTP, HTTPS, TCP, SSL, or Custom
- Amazon Auto Scaling & Elastic Load Balance can work together

Components of an Auto Scaling system

- Load balancer
- Solutions to measure the performance of current setup
- Scaling policy defining when to scale
- Resource provisioning policy
- Dynamic deployment template

Cloud-based Performance – Open solutions

- Multiple approaches are possible
- Shell
 - Linux utilities
 - Default
 - free –m

```
$ iostat -k
Linux 2.6.32-31-generic (oinas-laptop) 03/16/2012 _i686_ (2 CPU)

avg-cpu:  %user   %nice %system %iowait  %steal   %idle
           10.03    1.29    5.48    1.64    0.00   81.55

Device:            tps    kB_read/s    kB_wrtn/s    kB_read  kB_wrtn
sda                  1.71         6.42        13.82   31745551  68287580
```


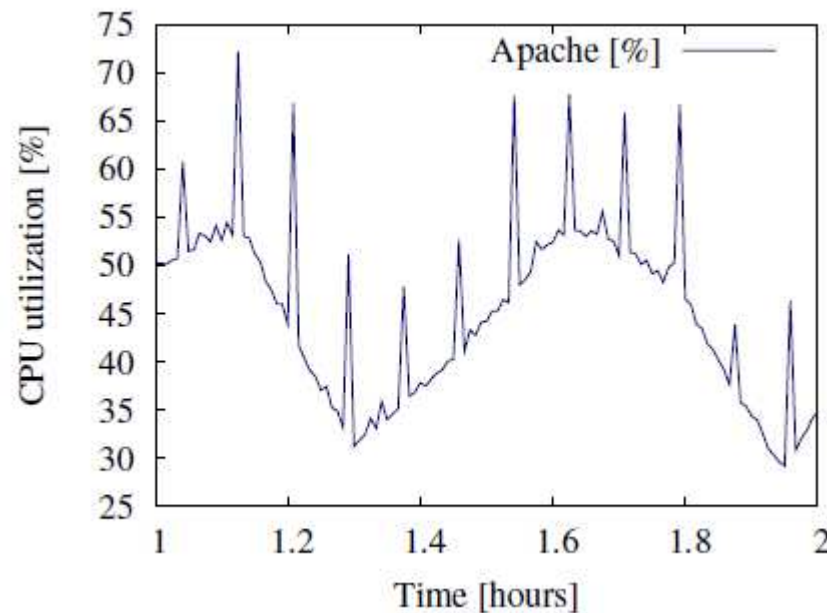
Cloud-based Performance - continued

- Tools (distributed)
 - Collectd
 - RRDtool
 - Generating visual performance graphs
 - Multicast communication
 - Does not impact system performance
 - Cacti
 - RRD
 - GUI
 - Performance decreases by 20%

RRDtool - round-robin database tool

Cloud-based Performance - continued

- Cacti
 - Spikes denote gathering performance metrics



Scaling Policy

- Time based
 - Already seen with Amazon Auto Scale
 - E.g. every Friday at 13:00:00 or Feb 15th 10 more servers for Estonian tax board
 - Good for On & Off! and Consistent traffic patterns
- Reactive
 - Threshold-based scaling policies
 - E.g. CPU utilization of all servers on average is >75% in last 5 min
 - Good for Fast Growth traffic pattern
- Predictive
 - AutoScaling based on predictive traffic
 - E.g. Predicting next min load by taking mean of last 5 min load
 - Good for Variable traffic pattern

Components of an Auto Scaling system

- Load balancer
- Solutions to measure the performance of current setup
- Scaling policy defining when to scale
- **Resource provisioning policy**
- **Dynamic deployment template**


Resource provisioning policy

- Simple resource provisioning policy
 - Resources estimation based on heuristic
 - E.g. suppose a node supports ~ 10 rps and current setup has 4 servers and load is 38 rps
 - Assume load increased or predicted to increase to 55 rps
 - So add 2 more servers
- May not be optimal or perfect solution, but sufficient for the immediate goals

Optimal Resource Provisioning for Auto-Scaling Enterprise Applications



- Cloud providers offer various instance types with different processing power and price
 - Can it be exploited in deciding the resource provisioning policy?
 - Makes the **policy to be aware of current deployment configuration**
- Another challenge: **Cloud providers charge the resource usage for fixed time periods**
 - E.g. Hourly prices of Amazon cloud
- Developed an LP based optimization model which considers both the issues [Srirama and Ostovar, CloudCom 2014]

Scaling enterprise application with the optimization model



Incoming load and scaling curves of Optimization model

Instance type usage curves of Optimization model



Scaling with Amazon AutoScale

[Srirama and Ostovar, CloudCom 2014]

Optimization Model

Intuition behind instance lifetime consideration


- Consider 2 instance types
 - Small instance(PW = 6r/s, Price = \$0.25/h),
 - Medium instance(PW = 12r/s, Price = \$0.4/h)

Load is 6r/s




Load increases to 12r/s => ?

Solution 1



Solution 2



- Saved cost with solution 2 : $0.46 - 0.44 = 0.02\$$
- So can we find this automatically?

Optimization Model

Some key definitions

- Region:
 - A task with its own independent characteristics
 - Each region can have its own capacity of instances
- Instance Type:
 - Each region can include multiple instance types
 - It is associated with processing power, price per period, capacity constraint, and configuration time
- Time bags:
 - Time interval where an instance is at a particular time
- Killing Cost:
 - Money lost when an instance is killed before it fills its paid period
- Retaining Cost:
 - The cost of the lived duration of the paid period

[Srirama and Ostovar, CloudCom 2014]

Optimization Model

- Cost Function:**

$$\begin{aligned}
 \text{Min } & \left(\sum_{i=1}^n \sum_{j=1}^m N_{r_i,t_j} * C_{r_i,t_j} + N_{r_i,t_j} * (CT_{r_i,t_j} * CTB_{r_i,t_j}) + \right. \\
 & \left. \sum_{i=1}^n \sum_{j=1}^m \sum_{k=1}^q S_{r_i,t_j,tb_k} * KC_{r_i,t_j,tb_k} + \right. \\
 & \left. \sum_{i=1}^n \sum_{j=1}^m \sum_{k=1}^q (X_{r_i,t_j,tb_k} - S_{r_i,t_j,tb_k}) * RC_{r_i,t_j,tb_k} \right)
 \end{aligned}$$

→ Cost of new instances
→ Configuration cost of new instances
→ Cost of killed instances
→ Cost of retained instances

- Constraints:**

$$\begin{aligned}
 \sum_{j=1}^m (N_{r,t_j} + (\sum_{k=1}^q X_{r,t_j,tb_k} - S_{r,t_j,tb_k})) * P_{r,t_j} &\geq W_r \longrightarrow \text{Workload constraint} \\
 \sum_{j=1}^m (N_{r,t_j} + (\sum_{k=1}^q X_{r,t_j,tb_k} - S_{r,t_j,tb_k})) &\leq CC_r \longrightarrow \text{Cloud capacity constraint} \\
 N_{t_r} + (\sum_{k=1}^q X_{t_r,tb_k} - S_{t_r,tb_k}) &\leq CCT_{t_r} \longrightarrow \text{Instance type capacity constraint} \\
 S_{tb_{r,t}} &\leq X_{tb_{r,t}} \longrightarrow \text{Shutdown constraint} \\
 N_{r,t} &\geq 0 \\
 S_{r,t} &\geq 0
 \end{aligned}$$

Application of the model

- Identify the scalable components in an enterprise application
- Scalable components are load tested on the planned cloud
 - To extract application specific parameters of the model
- Incoming load of each region is extracted and fed to the optimization model
 - Produces the ideal deployment configuration of the application

Evaluation of the optimization model

- The optimization model performs at least as good as Amazon AutoScale
 - Sometimes outperforms in efficiency and mostly in response times
 - Further optimizations with scaling policy can also save cost
- The model is generic and can be applied to any cloud
 - Which follows similar utility computing model
- It is also applicable to the systems which need to span across multiple clouds
- The latencies are also reasonable
 - The model could always find the optimal solution within decent amount of time

[Srirama and Ostovar, CloudCom 2014]

Components of an Auto Scaling system

- Load balancer
- Solutions to measure the performance of current setup
- Scaling policy defining when to scale
- Resource provisioning policy
- **Dynamic deployment template**

CloudML



- Developed in REMICS EU FP7 project
- Developed to tame cloud heterogeneity
- Domain-specific language (DSL) for modelling the provisioning and deployment at design-time
 - Nodes, artefacts and bindings can be defined
- Different means to manipulate CloudML models
 - Programmatically via Java API
 - Declaratively, via serialized model (JSON)
- Models@Runtime
 - Dynamic deployment of CloudML based mod

```
"nodeTypes": [  
  {  
    "id": "SmallGNUlinux",  
    "os": "GNUlinux",  
    "compute": [ 2, 4 ],  
    "memory": [ 2048, 4096 ],  
    "storage": [ 10240 ],  
    "location": "eu",  
    "provides": [  
      { "id": "SSHCapability" }  
    ]  
  }  
]  
  
"artefactsTypes": [  
  {  
    "id": "Docs",  
    "retrieval": "wget http://cloudml.org/apps/docs_configure; wget  
http://cloudml.org/apps/docs_deploy",  
    "configuration": "sudo docs_configure",  
    "deployment": "sudo docs_deploy",  
    "requires": [  
      { "id": "JettyCapability" },  
      { "id": "MongoDBCapability" }  
    ]  
  }  
  ...  
]
```

TOSCA

- Topology & Orchestration Specification of Cloud Application
- By OASIS
 - Sponsored by IBM, CA, Rackspace, RedHat, Huawei and Others
- Goal: cross cloud, cross tools orchestration of applications on the Cloud
- Node Type
- Relationship Type
- TOSCA Template

<https://cloudify.co/2015/07/21/what-is-TOSCA-cloud-application-orchestration-tutorial-cloudify.html>

Final Thoughts on AutoScaling

- AutoScaling can be dangerous
 - E.g. Distributed Denial of Service (DDoS) attack
 - Have min-max allocations
- Choose the right metrics
 - Stay with basic metrics
 - CPU, mem, I/O disk/net etc.
 - Review autoscaling strategy with metrics
- Choose your strategy
 - Scale up early and Scale down slowly
 - Don't apply the same strategy to all apps

This week in lab

- You work with load balancing

Next Lecture

- Introduction to MapReduce

References

- Amazon Web (Cloud) Services – documentation <http://aws.amazon.com/documentation/>
- Elastic Load balancing <http://aws.amazon.com/elasticloadbalancing/>
- Load balancing - algorithms <https://devcentral.f5.com/articles/intro-to-load-balancing-for-developers-ndash-the-algorithms#.UwSBpc4a5SM>
- Auto Scaling - Amazon Web Services <http://aws.amazon.com/autoscaling/>
- Cluet, M., Autoscaling Best Practices, <http://www.slideshare.net/lynxmanuk/autoscaling-best-practices>
- Ferry, N., Chauvel, F., Rossini, A., Morin, B., & Solberg, A. (2013). Managing multi-cloud systems with CloudMF. In Proceedings of the Second Nordic Symposium on Cloud Computing & Internet Technologies (pp. 38-45). ACM.
- M. Vasar, S. N. Srirama, M. Dumas: Framework for Monitoring and Testing Web Application Scalability on the Cloud, Nordic Symposium on Cloud Computing & Internet Technologies (NORDICLOUD 2012), August 20-24, 2012, pp. 53-60. ACM.
- S. N. Srirama, A. Ostovar: Optimal Resource Provisioning for Scaling Enterprise Applications on the Cloud, The 6th IEEE International Conference on Cloud Computing Technology and Science (CloudCom-2014), December 15-18, 2014, pp. 262-271. IEEE.
- S. N. Srirama, T. Iurii, J. Viil: Dynamic Deployment and Auto-scaling Enterprise Applications on the Heterogeneous Cloud, 9th IEEE International Conference on Cloud Computing (CLOUD 2016), June 27- July 2, 2016, pp. 927-932. IEEE.
- S. N. Srirama, A. Ostovar: Optimal Cloud Resource Provisioning for Auto-scaling Enterprise Applications, International Journal of Cloud Computing, ISSN: 2043-9997, 7(2):129-162, 2018. Inderscience.