



UNIVERSITY OF TARTU

INSTITUTE OF COMPUTER SCIENCE



Basics of Cloud Computing – Lecture 4

Introduction to MapReduce

Pelle Jakovits

Satish Srirama

Some material adapted from slides by Jimmy Lin, Web-Scale Information Processing Applications course, University of Waterloo (licensed under Creative Commons Attribution 3.0 License)

Outline

- Economics of Cloud Providers recap
- MapReduce model
- Hadoop MapReduce framework
- Hadoop Distributed File System (HDFS)
- Hadoop v2.0: YARN

Economics of Cloud Providers – Failures


- Cloud computing can greatly simplify large scale data processing by providing virtually unlimited access to storage and computing resources
- However, Cloud Computing brought a shift from highly reliable servers to commodity servers
- High number of servers means that failures are common.
 - Software must adapt to failures
- Solution: Replicate data and computations
 - Distributed File System & MapReduce
- MapReduce = functional programming meets distributed processing on steroids
 - Not a new idea - dates to the 50's

Functional Programming -> MapReduce

- Two important concepts in functional programming:
 - **Map:** Apply a user defined function on every element of the list
 - **Fold:** Apply a user defined aggregation function on a list to reduce it into a single value


Map

- Map is a higher-order function
 - **map** **f** **Lst**: $(a \rightarrow b) \rightarrow [a] \rightarrow [b]$
 - **f** - user defined function
 - **Lst** - a list of values
- Function **f** is applied to every element in the **input list**
- Result is a **new list**



Fold (FoldL)

- Fold is also a higher-order function
- **fold** **f** **x** **Lst**: $(a \rightarrow b \rightarrow a) \rightarrow a \rightarrow [b] \rightarrow a$
 - **f** - user defined function
 - **x** – initial accumulator value
 - **Lst** - a list of values



1. **Accumulator** is set to initial value **x**
2. Function **f** is applied to the first **list element** and the current value of the **accumulator**
3. Result is stored in the **accumulator**
4. Steps 2 and 3 are repeated for every following **item in the list**
5. The final value in the **accumulator** is returned as the result

Map/Fold in Action

- Simple map example:

`square x = x * x`

`map square [1,2,3,4,5] → [1,4,9,16,25]`

- Fold examples:

`fold (+) 0 [1,2,3,4,5] → 15`

`fold (*) 1 [1,2,3,4,5] → 120`

- Sum of squares:

`fold (+) 0 (map square [1,2,3,4,5])) ->`

`fold (+) 0 [1,4,9,16,25] -> 55`

Implicit parallelism

- In a purely functional setting, operations inside **map** can be performed independently and easily parallelised
 - We can partition the input list between multiple computers
 - We can apply the **map** operation on each partition separately
 - **Fold** provides a mechanism for combining **map** results back together
- If function **f** is *associative*, then we can also compute **fold** operations on different partitions independently

$$f(f(a, b), c) = f(a, f(b, c))$$

$$(a + b) + c = a + (b + c)$$

$$\text{fold } f \times [1,2,3,4] = f(\text{fold } f \times [1,2], \text{fold } f \times [3,4])$$

- This is the "*implicit parallelism*" of functional programming that MapReduce aims to exploit

Typical Large-Data Problem

- Map* • Iterate over a large number of records
- Extract something of interest from each
- Shuffle and sort intermediate results
- Aggregate intermediate results
- Generate final output *Reduce*

Key idea: provide a functional abstraction for these two operations – MapReduce

MapReduce

- Designed for processing very large-scale data (Petabytes)
- Deployed on a large cluster of servers
- Horizontally scalable - Add more servers to handle more data and speed up processing
- Data is stored in a distributed file system and replicated
- Parallelism is achieved by executing many Map and Reduce tasks concurrently
- **NB!** MapReduce **map** and **reduce** functions are not exactly the same as **map** and **fold** functions from functional programming!

MapReduce model

- Input is a list of Key and Value pairs: $[(k, v)]$
 - For example: (**LineNr**, **LineString**) when processing text files
- Programmers only have to specify two functions:
 - map** $(k, v) \rightarrow (k', v')^*$
 - reduce** $(k', [v']) \rightarrow (k'', v'')^*$
- The execution framework handles everything else:
 - Data partitioning, distribution, synchronization, fault recovery, etc.

Map function

- Map function is applied to every (Key, Value) pair in the input list
- Input to the user defined map functions is a single (Key, Value) pair
- Output is zero or more key and value pairs.
 - In functional programming, map function always had to return exactly one value.

$$\text{map } (k, v) \rightarrow (k', v')^*$$

Reduce function

- All KeyPairs produced in the map step are grouped by the keys and values are combined into a list
 - This happens between Map and Reduce stages
- Input to a reduce function is a **unique key** and a list of **values**: (Key, [Value])
- Reduce function is applied on the key and list of values
 - Typically an aggregation function is applied on the list
- Output is zero or more key and value pairs

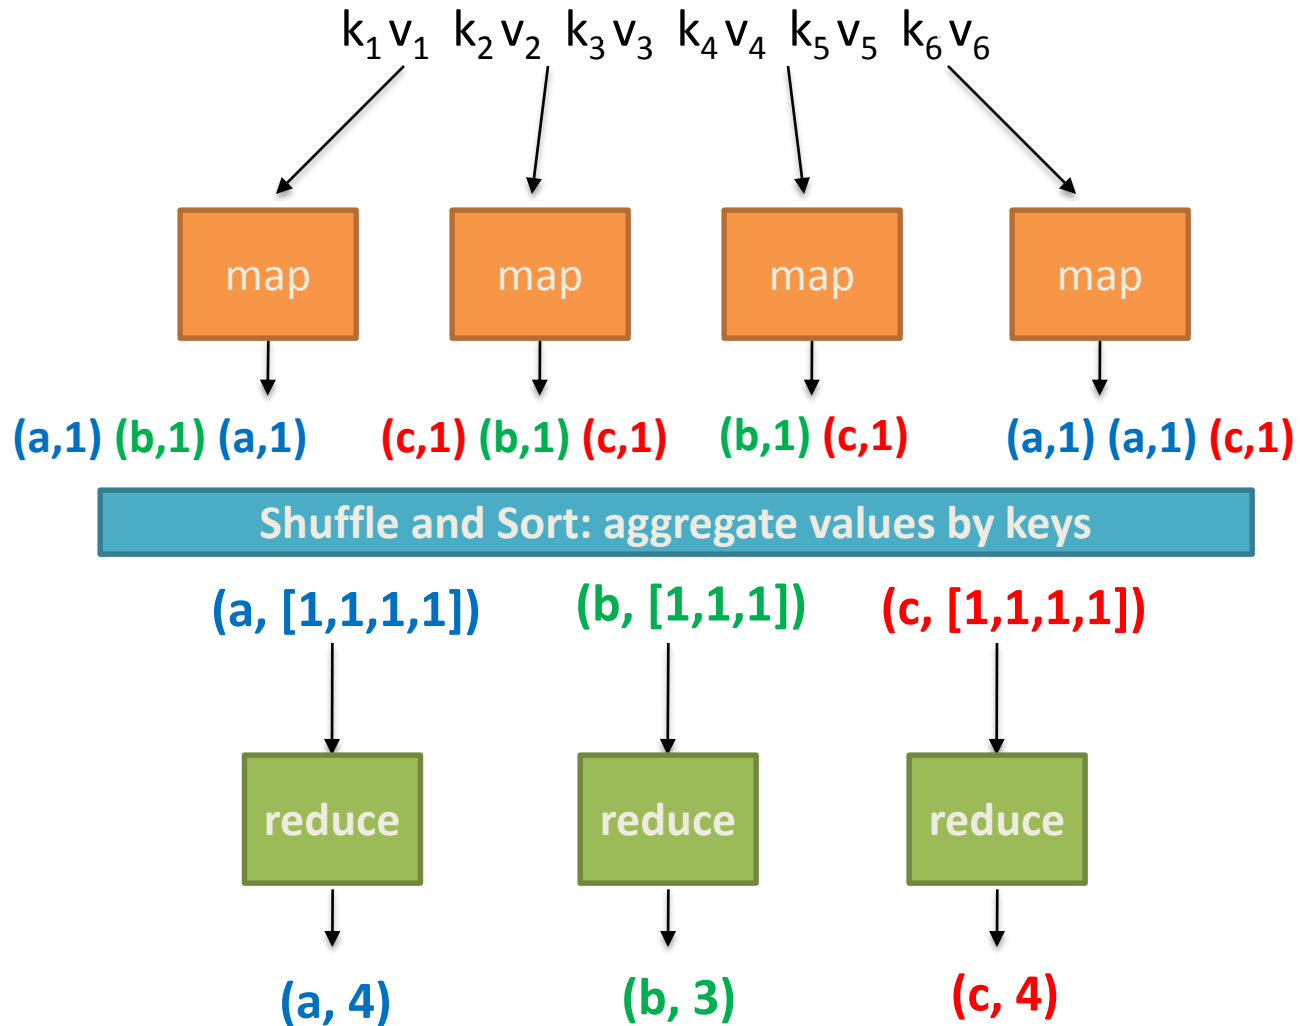
reduce (k', [v']) \rightarrow (k'', v'')*

Example: Word Count

```
Map(String lineNr, String line):  
  for each word w in line:  
    Emit(w, 1);
```

```
Reduce(String term, Iterator<Int> values):  
  int sum = 0;  
  for each v in values:  
    sum += v;  
  Emit(term, sum);
```

MapReduce



MapReduce

- Programmers specify two functions:
map $(k, v) \rightarrow (k', v')^*$
reduce $(k', [v']) \rightarrow (k'', v'')^*$
 - All values with the same key are sent to the same reducer
- The execution framework handles everything else...

What's “everything else”?

MapReduce “Runtime”

- Handles “data distribution”
 - Partition and replicate data
 - Moves processes to data
- Handles scheduling
 - Assigns workers to map and reduce tasks
- Handles synchronization
 - Gathers, sorts, and shuffles intermediate data
- Handles errors and faults
 - Detects worker failures and automatically restarts
- Handles speculative execution
 - Detects “slow” workers and re-executes work

Sounds simple, but many challenges!

MapReduce - extended

- Programmers specify two functions:
 - map** $(k, v) \rightarrow (k', v')^*$
 - reduce** $(k', [v']) \rightarrow (k'', v'')^*$
 - All values with the same key are reduced together
- The execution framework handles everything else...
- Users can change the behaviour of the framework by overwriting certain components, such as:
 - **Partitioner**
 - **Combiner**

Partitioner

- Controls how the key space is distributed into different partitions.
`Partition(key, number of partitions) → partition for key`
- A simple hash of the key can be used, e.g.:
`Partition(Key, partitions) = hashFun(Key) mod partitions`
- Programmers can overwrite partitioner function to force specific keys to be located in the same partition
 - Example: **Global sorting order**
 - Partition keys that start with **0** into partition **0**, **1** into **1**, **2** into **2**, ..., **9** into **9**.
 - Makes sure that every previous partition has smaller keys than the next one.

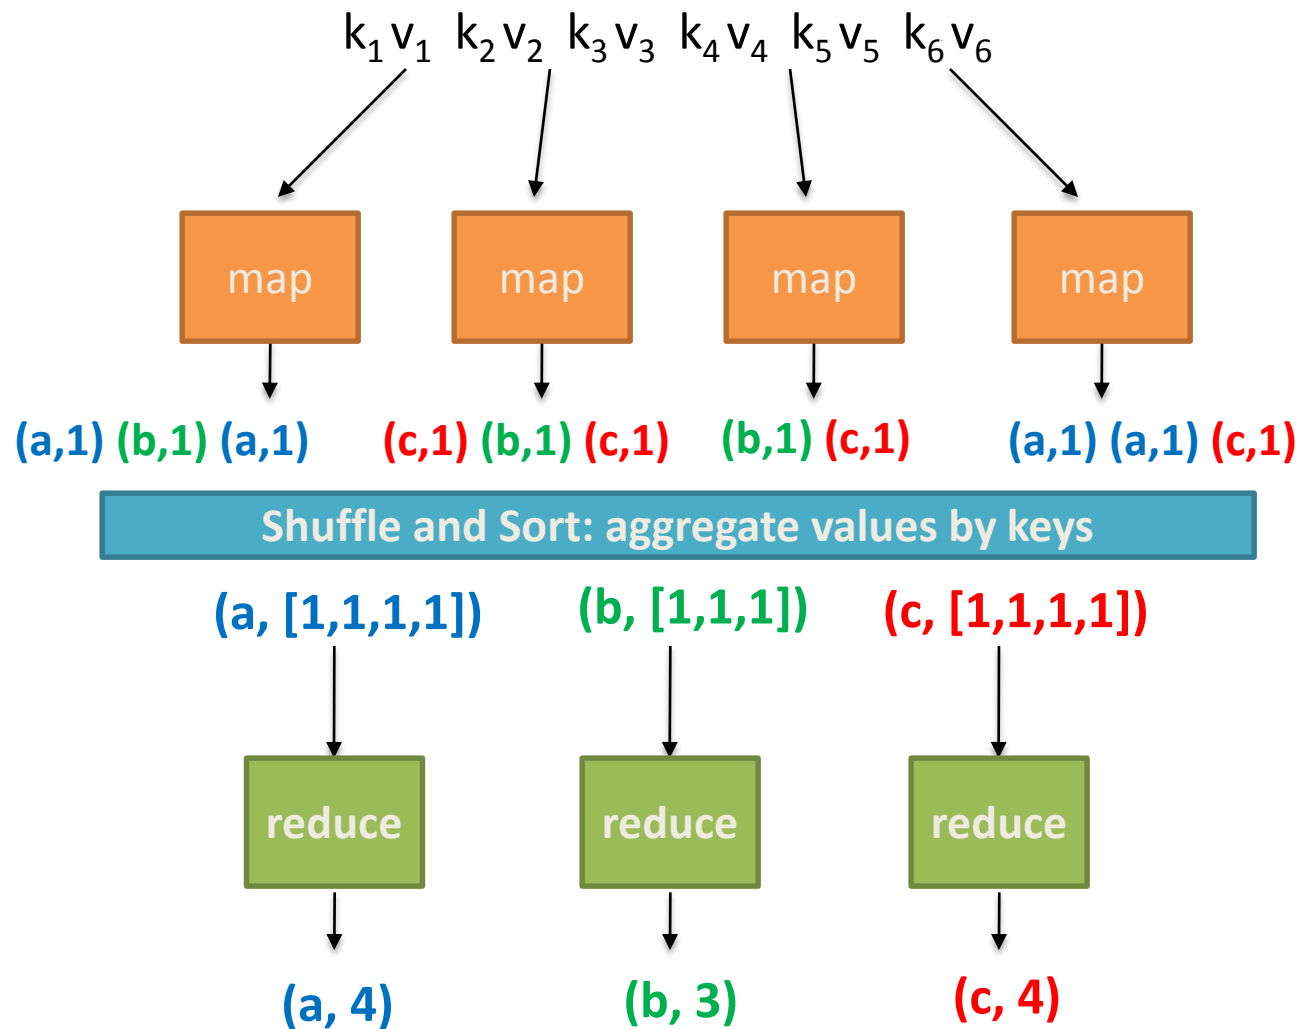
Combiner

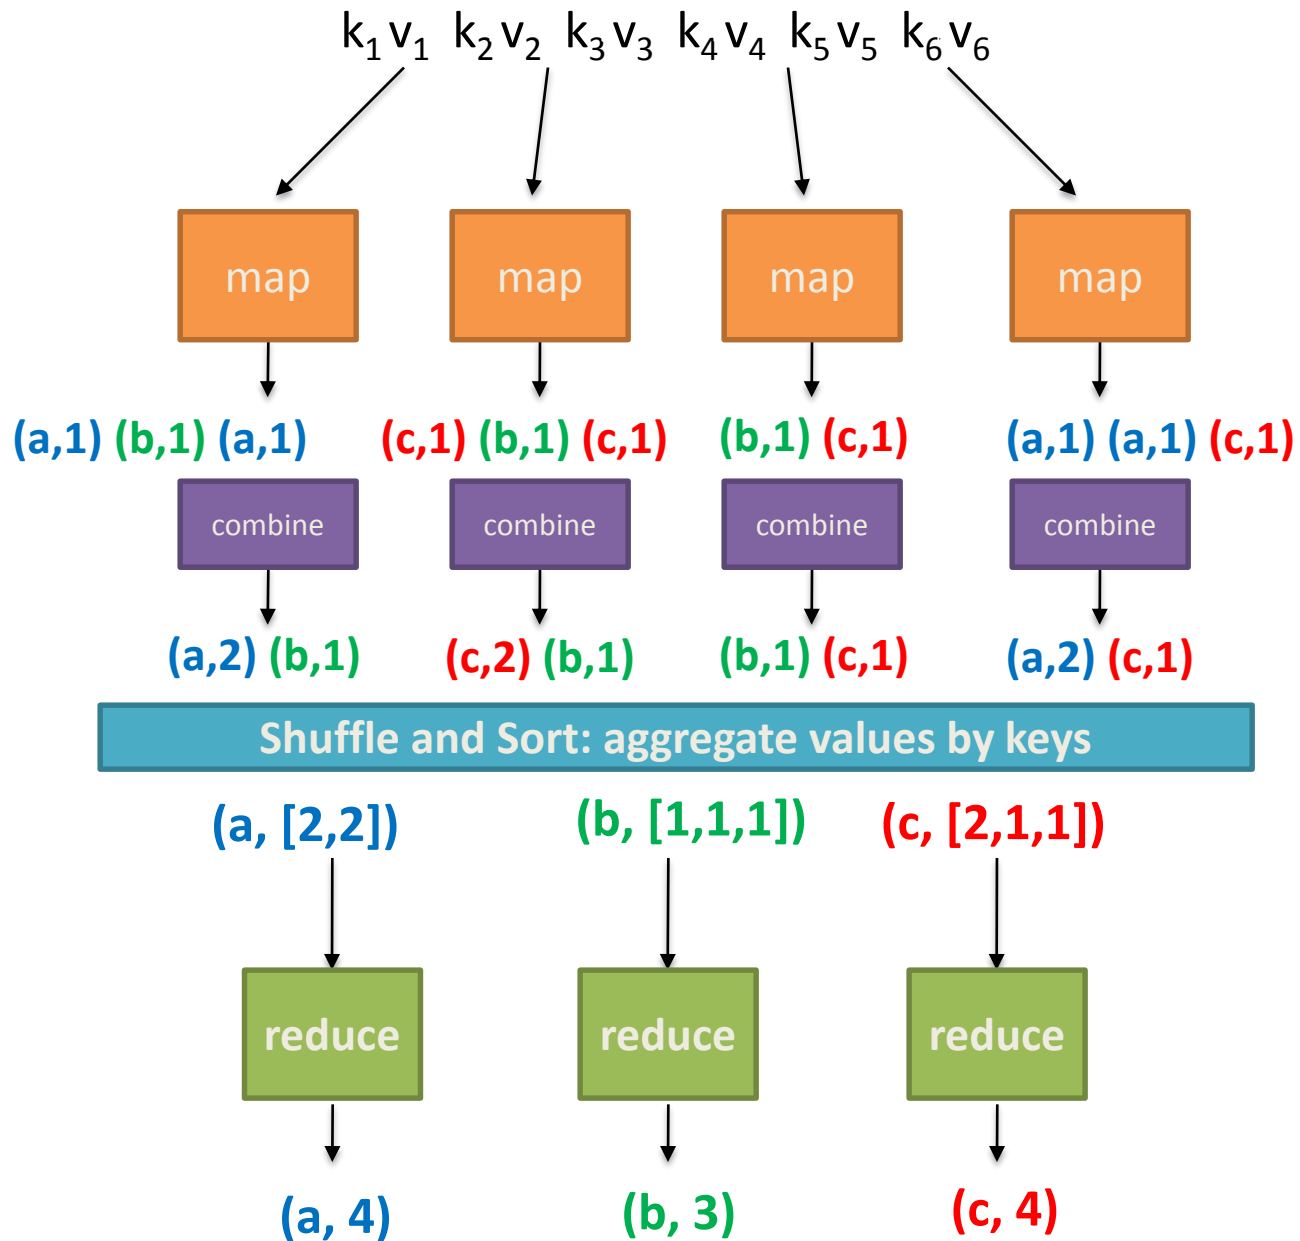
- In many cases, Map stage produces too many KeyValue pairs.
 - WordCount produces a KeyValue for every single word.
- Combine can be used as a mini-reducer stage that runs in memory after the map phase

combine $(k', [v']) \rightarrow (k'', v'')^*$

- Used as an optimization to reduce network traffic
- Often the Reduce function can be used directly as a combiner
 - But not always!
 - Why?

WordCount without Combiner

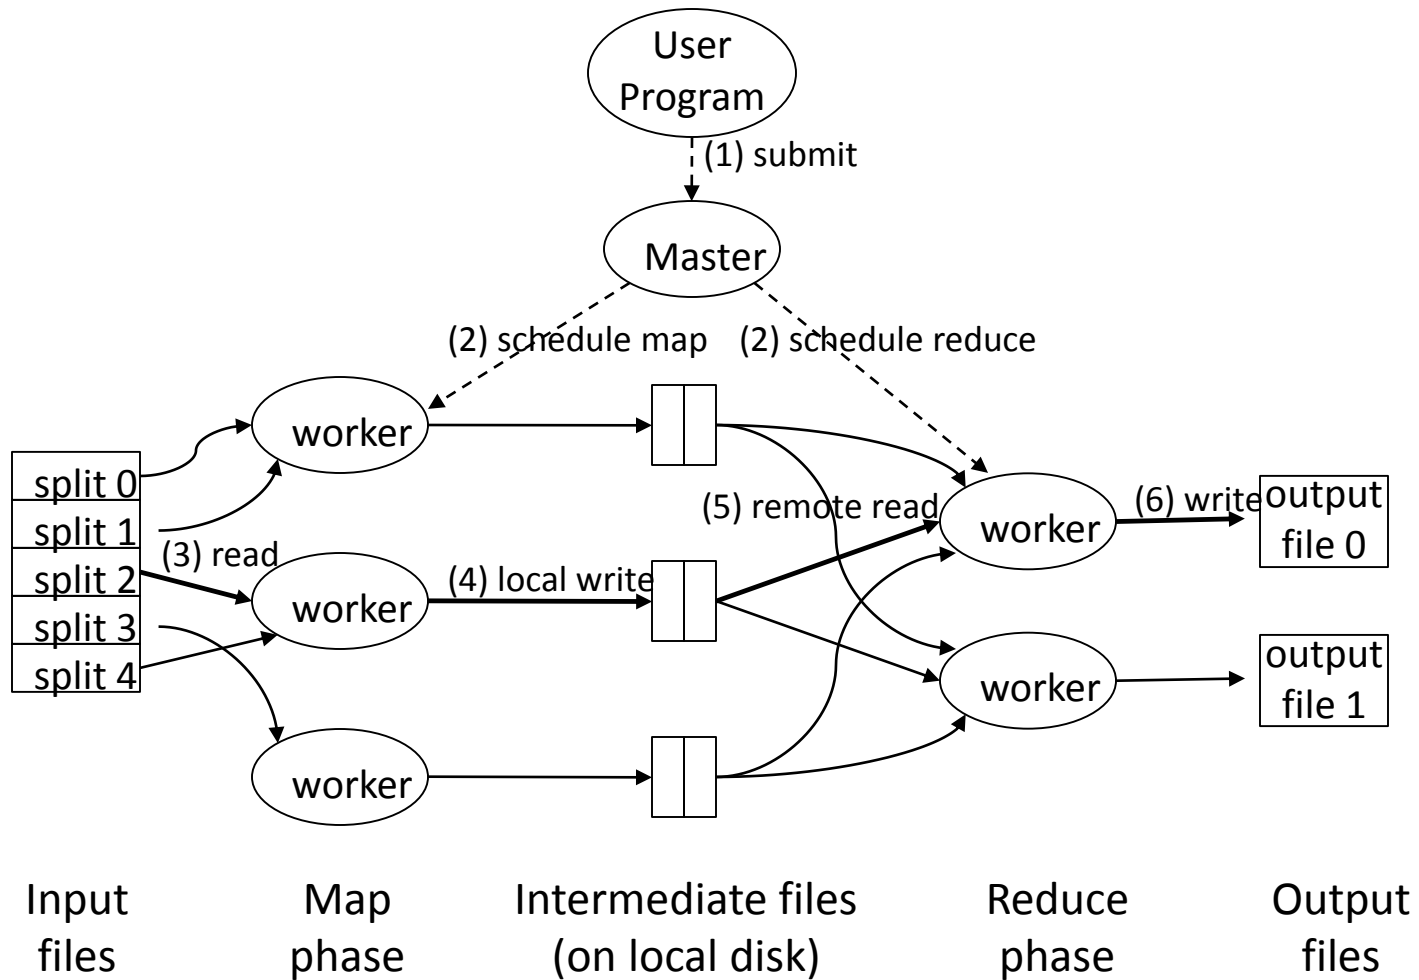




Two more important details...

- Barrier between map and reduce phases
 - Reduce tasks wait until map tasks are finished
 - But we can begin copying intermediate data earlier
- Data is ordered by keys before reduce operation is performed
 - But no enforced ordering *across* reducers
 - Unless using a specially designed partitioner!

MapReduce Overall Architecture



Adapted from (Dean and Ghemawat, OSDI 2004)

MapReduce can refer to...


- The programming model
- The execution framework (aka “runtime”)
- The specific implementation

Usage is usually understandable from context!

MapReduce Implementations

- Google has a proprietary implementation in C++
 - Bindings in Java, Python
- Hadoop is an open-source implementation in Java
 - Development led by Yahoo, used in production
 - Now an Apache project
 - Rapidly expanding software ecosystem
- Lots of custom implementations
 - For GPUs, cell processors, etc.
 - MapReduce as database query engine (CouchDB)

Cloud Computing Storage: How to move data to the workers?



Managing Peta-scale data

- Network bandwidth is limited
- Not enough RAM to hold all the data in memory
- Disk access is slow, but disk throughput is reasonable
- Don't move data to workers... move workers to the data!
 - Store data on the local disks of nodes in the cluster
 - Start up the workers on the node where the data is locally stored
- A distributed file system is the answer
 - GFS (Google File System) for Google's MapReduce
 - HDFS (Hadoop Distributed File System) for Hadoop

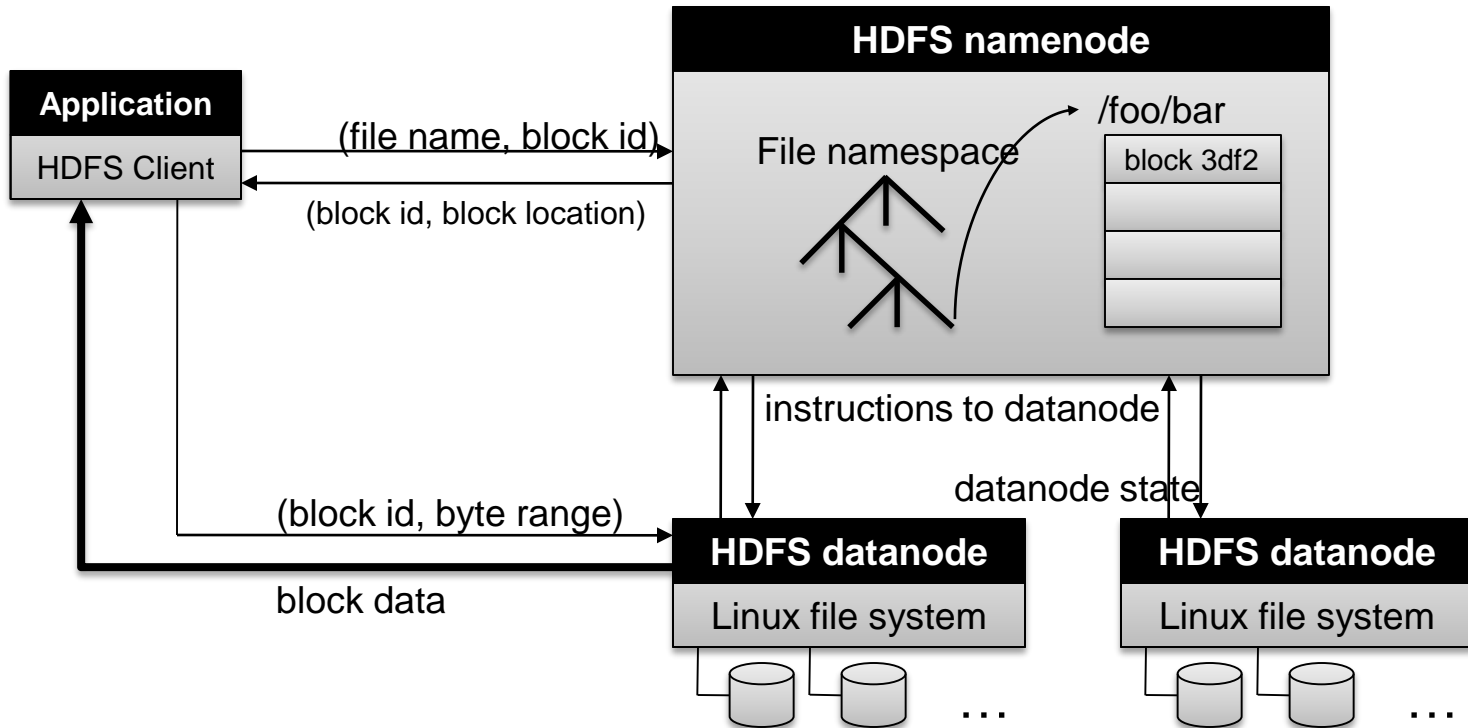
Distributed File System - Assumptions

- Choose cheaper commodity hardware over “exotic” hardware
 - Scale “out”, not “up”
- High component failure rates
 - Inexpensive commodity components fail all the time
- “Modest” number of huge files
 - Multi-gigabyte files are common, if not encouraged
- Files are write-once, mostly appended to
 - Perhaps concurrently
- Large streaming reads over random access
 - High sustained throughput over low latency

Distributed File System - Design Decisions

- Files stored as chunks
 - Fixed size (64MB)
- Reliability through replication
 - Each chunk replicated across 3+ nodes
- Single master to coordinate access, keep metadata
 - Simple centralized management
- No data caching
 - Little benefit due to large datasets, streaming reads
- Simplify the API
 - Push some of the issues onto the client (e.g., data layout)
- Hadoop Distributed File System (HDFS) is the implementation of Google File System (GFS)

HDFS Architecture



Namenode Responsibilities

- Managing the file system namespace:
 - Holds file/directory structure, metadata, file-to-block mapping, access permissions, etc.
- Coordinating file operations:
 - Directs clients to datanodes for reads and writes
 - No data is moved through the namenode
- Maintaining overall health:
 - Periodic communication with the datanodes
 - Block re-replication and rebalancing
 - Garbage collection

Hadoop Processing Model

- Create or allocate a cluster
- Put data onto the file system
 - Data is split into blocks
 - Replicated and stored in the cluster
- Run your job
 - Copy Map code to the allocated nodes
 - Move computation to data, not data to computation
 - Gather output of Map, sort and partition on key
 - Run Reduce tasks
- Results are stored in the HDFS

MapReduce Terminology

- **Job** – A “full program” - an execution of a Mapper and Reducer across a data set
- **Task** – An execution of a Mapper or a Reducer on a data chunk
 - Also known as Task-In-Progress (TIP)
- **Task Attempt** – A particular instance of an attempt to execute a task on a machine

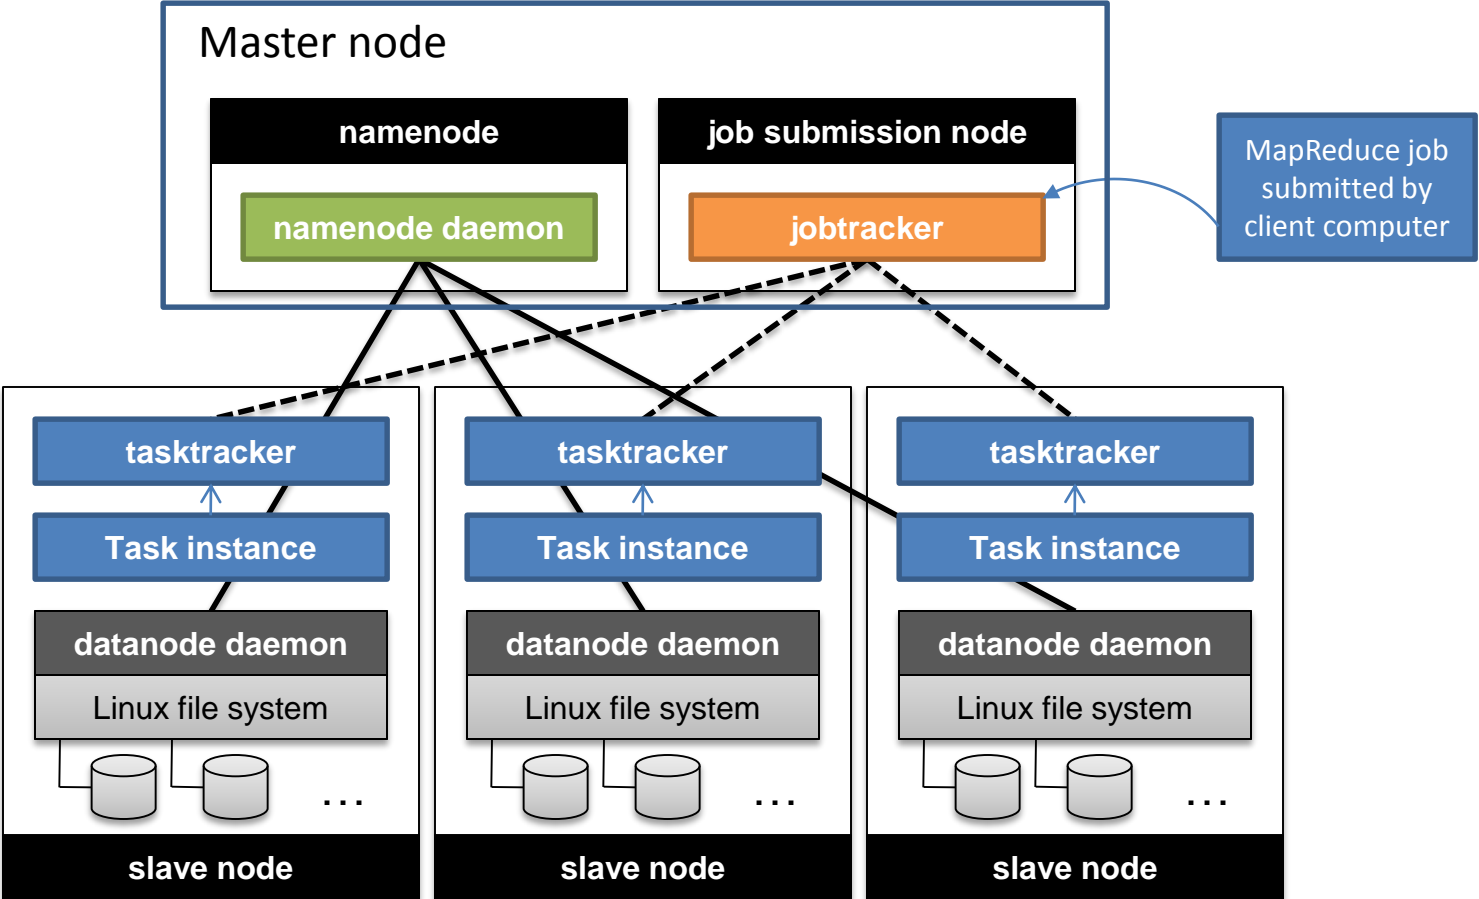
Terminology example

- Running “Word Count” across 20 files is one **job**
- 20 files to be mapped imply 20 map **tasks**
+ some number of reduce **tasks**
- At least 20 **map task attempts** will be performed
 - more if a machine crashes or slow, etc.

Task Attempts

- A particular task will be attempted at least once, possibly more times if it crashes
 - If the same input causes crashes over and over, that input will eventually be abandoned
- Multiple attempts at one task may occur in parallel with speculative execution turned on

Hadoop MapReduce Architecture : High Level



MapReduce Summary

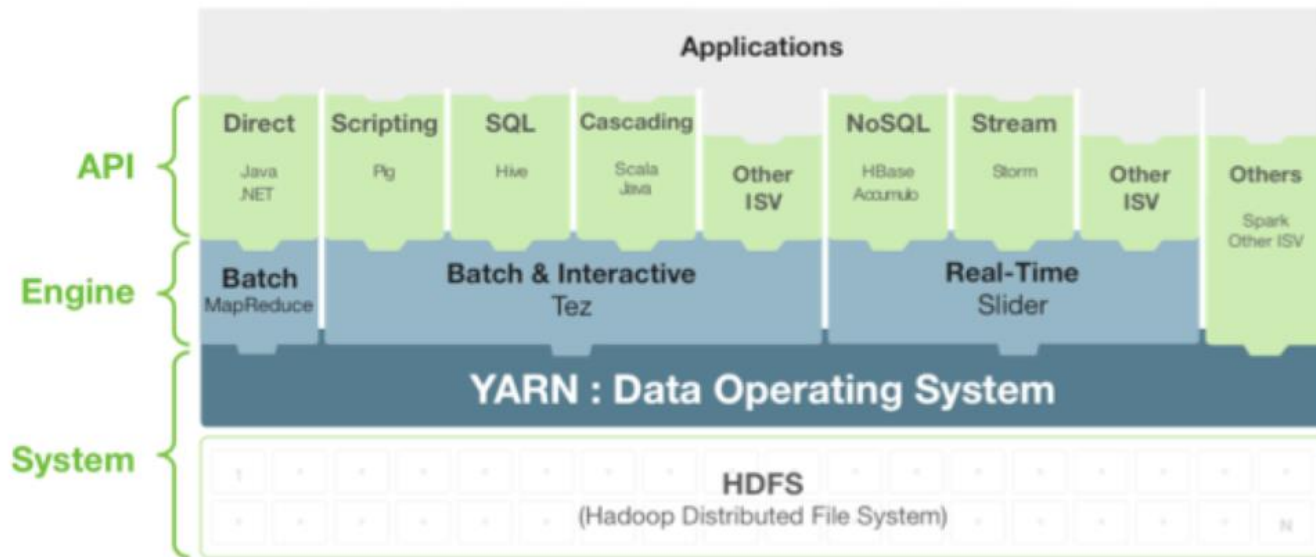
- Simple, but powerful programming model
- Scales to handle petabyte+ workloads
 - Google: six hours and two minutes to sort 1PB (10 trillion 100-byte records) on 4,000 computers
 - Yahoo!: 16.25 hours to sort 1PB on 3,800 computers
- Incremental performance improvement with more nodes
- Seamlessly handles failures, but possibly with performance penalties

Limitations with MapReduce V1

- Master node has too many responsibilities!
- Scalability issues
 - Maximum Cluster Size – 4000 Nodes
 - Maximum Concurrent Tasks – 40000
- Coarse synchronization in Job Tracker
 - Single point of failure
 - Failure kills all queued and running jobs
- Jobs need to be resubmitted by users
 - Restart is very tricky due to complex state
- Problems with resource utilization


MapReduce NextGen aka YARN aka MRv2

- New architecture introduced in hadoop-0.23
- Divides two major functions of the JobTracker
 - Resource management and job life-cycle management are divided into separate components
- An application is either a single job in the sense of classic MapReduce jobs or a DAG of such jobs



YARN Architecture

- **ResourceManager:**
 - Arbitrates resources among all the applications in the system
 - Has two main components: Scheduler and ApplicationsManager
- **NodeManager:**
 - Per-machine slave
 - Responsible for launching the applications' containers, monitoring their resource usage
- **ApplicationMaster:**
 - Negotiate appropriate resource containers from the Scheduler, tracking their status and monitoring for progress
- **Container:**
 - Unit of allocation incorporating resource elements such as memory, cpu, disk, network etc.
 - To execute a specific task of the application
 - Similar to map/reduce slots in MRv1



Execution Sequence with YARN

- A **client** program submits the application
- **ResourceManager** allocates a specified container to start the **ApplicationMaster**
- **ApplicationMaster**, on boot-up, registers with **ResourceManager**
- **ApplicationMaster** negotiates with **ResourceManager** for appropriate resource containers
- On successful container allocations, **ApplicationMaster** contacts **NodeManager** to launch the container
- Application code is executed within the container, and then **ApplicationMaster** is responded with the execution status
- During execution, the client communicates directly with **ApplicationMaster** or **ResourceManager** to get status, progress updates etc.
- Once the application is complete, **ApplicationMaster** unregisters with **ResourceManager** and shuts down, freeing its own container process

Next Lab

- Set up IDE for creating Hadoop MapReduce applications
 - Run Hadoop MapReduce code in your computer without installing Hadoop
- Try out MapReduce WordCount example
- Improve the WordCount example

Next Lecture

- We will take a look at different MapReduce algorithms
- Learn how to design MapReduce applications

References

- Hadoop wiki <http://wiki.apache.org/hadoop/>
- Data-Intensive Text Processing with MapReduce
 - Authors: Jimmy Lin and Chris Dyer
 - <http://lntool.github.io/MapReduceAlgorithms/MapReduce-book-final.pdf>
- White, Tom. Hadoop: the definitive guide. O'Reilly, 2012.
- J. Dean and S. Ghemawat, “MapReduce: Simplified Data Processing on Large Clusters”, OSDI'04: Sixth Symposium on Operating System Design and Implementation, Dec, 2004.
- Apache Hadoop YARN architecture - <https://hadoop.apache.org/docs/current/hadoop-yarn/hadoop-yarn-site/YARN.html>