



UNIVERSITY OF TARTU

INSTITUTE OF COMPUTER SCIENCE



Basics of Cloud Computing – Lecture 7

More AWS & Serverless Computing

Pelle Jakovits

Satish Srirama



Mobile & Cloud Lab



Outline

- Overview of Amazon Web Services
- Serverless computing
 - FaaS Model
 - Apache OpenWhisk
 - Advantages and disadvantages

Cloud Providers and Services we have discussed

- Amazon Web Services
 - **Compute:** EC2
 - **Storage:** S3, EBS
 - **Scaling:** Elastic Load Balancing, Auto Scale, CloudWatch
- Eucalyptus
- OpenStack
- Management providers
 - AWS Management Console
 - OpenStack Horizon
 - RightScale
- PaaS
 - Google AppEngine
 - Windows Azure
 - Elastic MapReduce

AWS services we will discuss

- Management Console
- Identity and Access Management
- CloudFormation
- Data Services
- Data Pipelines
- Data migration services

AWS Management Console

- Hope some of you have started using Amazon accounts
- You can manage your complete Amazon account with management console
 - AMI Management
 - Instance Management
 - Security Group Management
 - Elastic IP Management
 - Elastic Block Store
 - Key Pair management
 - etc.
- Have different pages for different services

AWS services

Find Services

You can enter names, keywords or acronyms.

 Example: Relational Database Service, database, RDS

▼ Recently visited services



[ELB](#)



[Data Pipeline](#)



[OpsWorks](#)



[CloudFormation](#)




[SWF](#)

▼ All services



Compute

[EC2](#)
[Lightsail](#) 
[ECR](#)
[ECS](#)
[EKS](#)
[Lambda](#)
[Batch](#)
[Elastic Beanstalk](#)
[Serverless Application Repository](#)



Storage

[S3](#)
[EFS](#)
[FSx](#)
[S3 Glacier](#)



Developer Tools

[CodeStar](#)
[CodeCommit](#)
[CodeBuild](#)
[CodeDeploy](#)
[CodePipeline](#)
[Cloud9](#)
[X-Ray](#)



Robotics

[AWS RoboMaker](#)



Blockchain

[Amazon Managed Blockchain](#)



[Satellite](#)



Machine Learning

[Amazon SageMaker](#)
[Amazon Comprehend](#)
[AWS DeepLens](#)
[Amazon Lex](#)
[Machine Learning](#)
[Amazon Polly](#)
[Rekognition](#)
[Amazon Transcribe](#)
[Amazon Translate](#)
[Amazon Personalize](#)
[Amazon Forecast](#)
[Amazon Textract](#)



Analytics

[Athena](#)



Mobile

[AWS Amplify](#)
[Mobile Hub](#)
[AWS AppSync](#)
[Device Farm](#)



AR & VR


[Amazon Sumerian](#)



Application Integration

[Step Functions](#)
[Amazon MQ](#)
[Simple Notification Service](#)
[Simple Queue Service](#)
[SWF](#)

AWS EC2 DashBoard



Services ▾ Resource Groups ▾ ⚙

EC2 Dashboard

Events

Tags

Reports

Limits

INSTANCES

Instances

Launch Templates

Spot Requests

Reserved Instances

Dedicated Hosts

Scheduled Instances

Capacity Reservations

IMAGES

AMIs

Bundle Tasks

ELASTIC BLOCK STORE

Volumes

Snapshots

Lifecycle Manager

NETWORK & SECURITY

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

LOAD BALANCING

Load Balancers

Target Groups

Resources

You are using the following Amazon EC2 resources in the EU West (Ireland) region:

0 Running Instances

0 Dedicated Hosts

0 Volumes

17 Key Pairs

0 Placement Groups

0 Elastic IPs

1 Snapshots

0 Load Balancers

32 Security Groups

Learn more about the latest in AWS Compute from AWS re:Invent by viewing the [EC2 Videos](#).

Create Instance

To start using Amazon EC2 you will want to launch a virtual server, known as an Amazon EC2 instance.

Launch Instance ▾

Note: Your instances will launch in the EU West (Ireland) region

Service Health

Service Status:

✓ EU West (Ireland):

Availability Zone Status:

✓ eu-west-1a:
Availability zone is operating normally

✓ eu-west-1b:
Availability zone is operating normally

✓ eu-west-1c:
Availability zone is operating normally

[Service Health Dashboard](#)

Scheduled Events

EU West (Ireland):

No events

AWS Identity and Access Management (IAM)

- How can an enterprise or group of users use one credit card?
- Manage IAM users
 - Create new users and manage them
 - Create groups
- Manage credentials
 - Create and assign temporary security credentials
- Manage permissions
 - Creating policies for specific services and users
 - Can use very fine-grained granularity

IAM policy example

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Action": [
      "ec2:DescribeInstances", "ec2:DescribeImages",
      "ec2:DescribeKeyPairs", "ec2:DescribeVpcs",
      "ec2:DescribeSubnets", "ec2:DescribeSecurityGroups"
    ],
    "Resource": "*"
  },
  {
    "Effect": "Allow",
    "Action": "ec2:RunInstances",
    "Resource": [
      "arn:aws:ec2:sa-east-1:111122223333:network-interface/*",
      "arn:aws:ec2:sa-east-1:111122223333:volume/*",
      "arn:aws:ec2:sa-east-1:111122223333:key-pair/*",
      "arn:aws:ec2:sa-east-1:111122223333:security-group/*",
      "arn:aws:ec2:sa-east-1:111122223333:subnet/subnet-1a2b3c4d"
    ]
  }
],
}
```

IAM policy example

```
{
  "Effect": "Allow",
  "Action": "ec2:RunInstances",
  "Resource": [
    "arn:aws:ec2:sa-east-1:111122223333:instance/*"
  ],
  "Condition": {
    "StringEquals": {
      "ec2:InstanceType": "m1.small"
    }
  }
},
{
  "Effect": "Allow",
  "Action": "ec2:RunInstances",
  "Resource": [
    "arn:aws:ec2:sa-east-1::image/ami-*"
  ],
  "Condition": {
    "StringEquals": {
      "ec2:Owner": "amazon"
    }
  }
}
]
```

AWS CloudFormation

- Provides an easy way to create and manage a collection of related AWS resources, provisioning and updating them in an orderly and predictable fashion
- It is based on templates model
 - Templates describe the AWS resources, the associated dependencies, and runtime parameters to run an app.
 - The templates describe stacks, which are set of software and hardware resources.
 - Something similar to CloudML and RightScale server templates
- Hides several details
 - How the AWS services need to be provisioned
 - Subtleties of how to make those dependencies work.

Resources:

Ec2Instance:

Type: 'AWS::EC2::Instance'

Properties:

SecurityGroups:

- !Ref **InstanceSecurityGroup**
- MyExistingSecurityGroup

KeyName: **mykey**

InstanceType: **t2.micro**

ImageId: **ami-7a11e213**

InstanceSecurityGroup:

Type: 'AWS::EC2::SecurityGroup'

Properties:

GroupDescription: Enable SSH access via port 22

SecurityGroupIngress:

- IpProtocol: tcp

FromPort: '**22**'

ToPort: '**22**'

CidrIp: 0.0.0.0/0

AWS CloudFormation

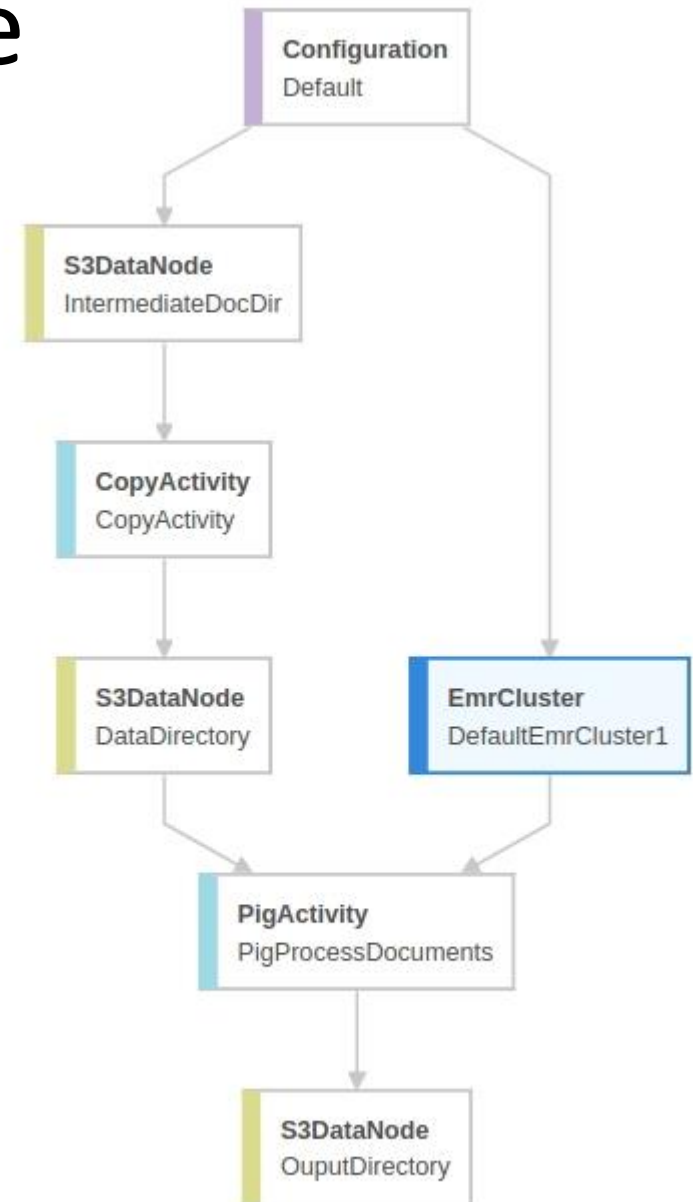
- Amazon provides several pre-built templates to start common apps, such as:
 - WordPress (blog)
 - LAMP stack (Linux, Apache, MySQL, and PHP)
 - Gollum (wiki used by GitHub)
- There is no additional charge for AWS CloudFormation.
- You pay for the utilized AWS resources (e.g. EC2 instances, Elastic Load Balancers, etc.)
- <http://aws.amazon.com/cloudformation/>

AWS Data Services

- Amazon Relational Database Service
 - Scalable and re-sizable SQL DB service
 - Supports most of the familiar database engines
 - MySQL, PostgreSQL, Oracle, Microsoft SQL Server
- Amazon Aurora
 - High Performance SQL service (MySQL and PostgreSQL compatible)
 - Distributed, fault-tolerant, self-healing storage
- NoSQL databases
 - DocumentDB - Managed Document DB (MongoDB compatible)
 - DynamoDB – Managed NoSQL database
 - Neptune – Managed Graph Database

AWS Data Pipeline

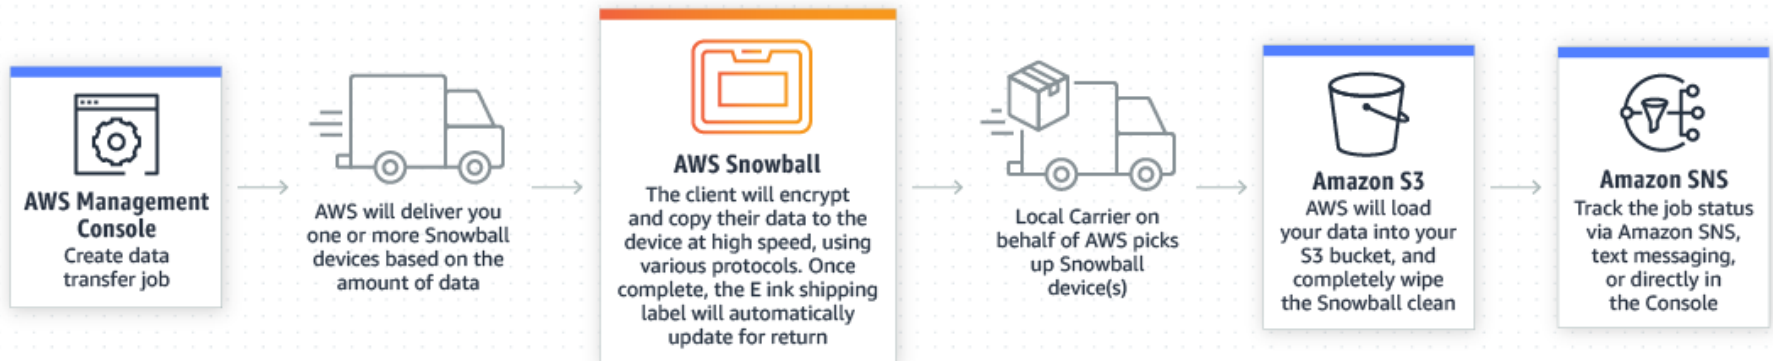
- Service for orchestrating data movement and processing tasks inside AWS
- **DataNode** - Location of the data source or destination. (SqlDataNode, S3DataNode, ...)
- **Activity** - Operation to perform on data (CopyActivity, EmrActivity, ...)
- **Schedule** - When data pipelines activities are initiated (On-demand, CRON, ...)
- **Precondition** - Conditions for when pipeline tasks can be executed
- **Resource** - EC2 resources or other AWS services
Activities depend on
- Open Source alternative: **Apache NiFi**



AWS Snowball



- Device for secure and physical Data Migration
- Storage capacity: 50 – 80TB
- Migrate Big Data: analytics data, genomics data, video libraries, image repositories, backups, etc.




AWS Snowball Edge

- Device for Data Migration together with Onboard pre-processing capability.
- Onboard compute service options:
 - AWS Lambda, EC2 AMIs
 - 26 or 52 vCPUs
 - **Optional GPU** for machine learning and real time video analysis
- Storage capacity: 100TB
- Clustering: Up to 20 nodes



AWS Snowmobile

- Petabyte- to Exabyte-scale data transfer service for moving extremely large amounts of data to AWS
- 100 PB of Data => as much as 1,250 AWS Snowball devices
- Pricing:
 - \$0.005/GB per month
 - \$5243/PB per month



Other AWS Services

- **Analytics** (EMR, Athena serverless queries, Kinesis video analytics)
- **Machine Learning** (DeepLens, SageMaker, Lex)
- **Application Integration** (Message Queue, Step Functions)
- **Internet of Things**
- **Media Services** (Transcoding pipelines, stream management)
- **Networking** (VPC, Domains)
- **Augmented & Virtual Reality**
- **RoboMaker** (Cloud services for Robot Operating System)
- **Security** (User management, Firewalls, Data privacy tools)
- **Blockchain**

A BIT MORE ON SERVERLESS COMPUTING

What is Serverless?



Abstraction
of servers



Event-driven/
instant scale



Sub-second
billing

Serverless computing - continued

- Newer workloads are a better fit for event driven programming
 - Execute application logic in response to database triggers
 - Execute app logic in response to sensor data
 - Execute app logic in response to scheduled tasks etc.
- Applications are charged by compute time (millisecond) rather than by reserved resources
- Greater linkage between cloud resources used and business operations executed
- Serverless in a nutshell
 - Event-action platforms to execute code in response to events

Current platforms for serverless



AWS Lambda

OpenLambda



Azure Functions



Red-Hat



Google Functions



Kubernetes

FaaS in Public Clouds

- **AWS Lambda**

- Run code in AWS without managing infrastructure or software
- Java, Go, PowerShell, Node.js, C#, Python, and Ruby code
- Pricing is based on number of **requests** and **GB-Sec "Memory-Duration"**
- Free: 1M **requests** a month. After: \$0.20 per 1M
- Free: 400,000 **GB-Sec**. After: \$0.000017 per 1 **GB-Sec**

- **IBM BlueMix Cloud Function**


- Based on OpenWhisk - Open Source Serverless cloud platform
- Event, trigger & rule based execution
- Supports any language*
- Free: 400,000 **GB-Sec**. After: \$0.000017 per 1 **GB-Sec**

Apache OpenWhisk

- Initiated by IBM but now an Apache project
- Open source cloud platform
- Serverless deployment and operations model
- Optimal utilization and granular pricing
- Scales on a per-request basis
- Supports JS, Swift, Python, Java, Docker

Triggers, actions, rules (and packages)

- Services or data sources define the events they emit as **triggers**
- Developers associate the **actions** to handle the events via **rules**
- **Packages** are a shared collection of Triggers and Actions



Triggers and Rules

- Trigger examples
 - changes to database records
 - IoT sensor readings that exceed a certain temperature
 - new code commits to a GitHub repository
 - simple HTTP requests from web or mobile apps
- Rule is an association of a trigger to an action
 - Many to many mapping

Actions

- They can be
 - small snippets of JavaScript or Swift code
 - custom binary code embedded in a Docker container
- Instantly deployed and executed whenever a trigger fires
- It is also possible to directly invoke an action by using the OpenWhisk API, CLI, or iOS SDK
- A set of actions can be chained

CLI – Command line interface

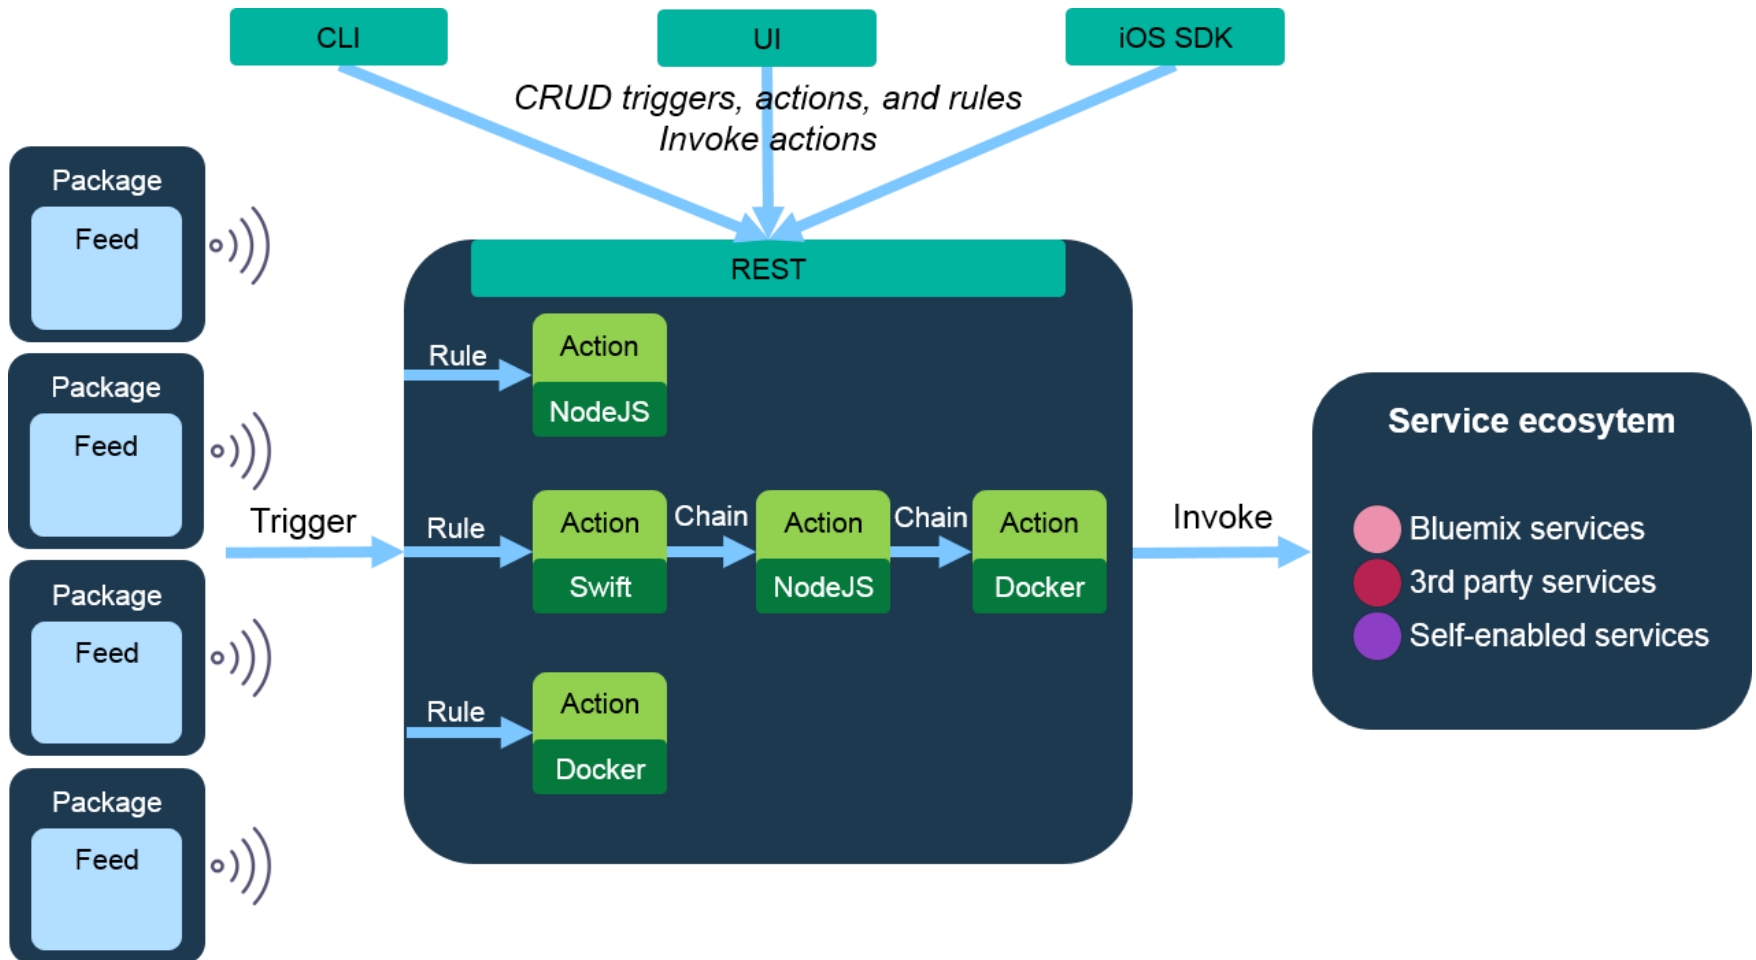
Actions - continued

Hello world as an OpenWhisk action.

```
def myFunction(args):  
    name = args['name']  
    greeting = "Hello " + name + '!'  
    return {"greeting": greeting}
```

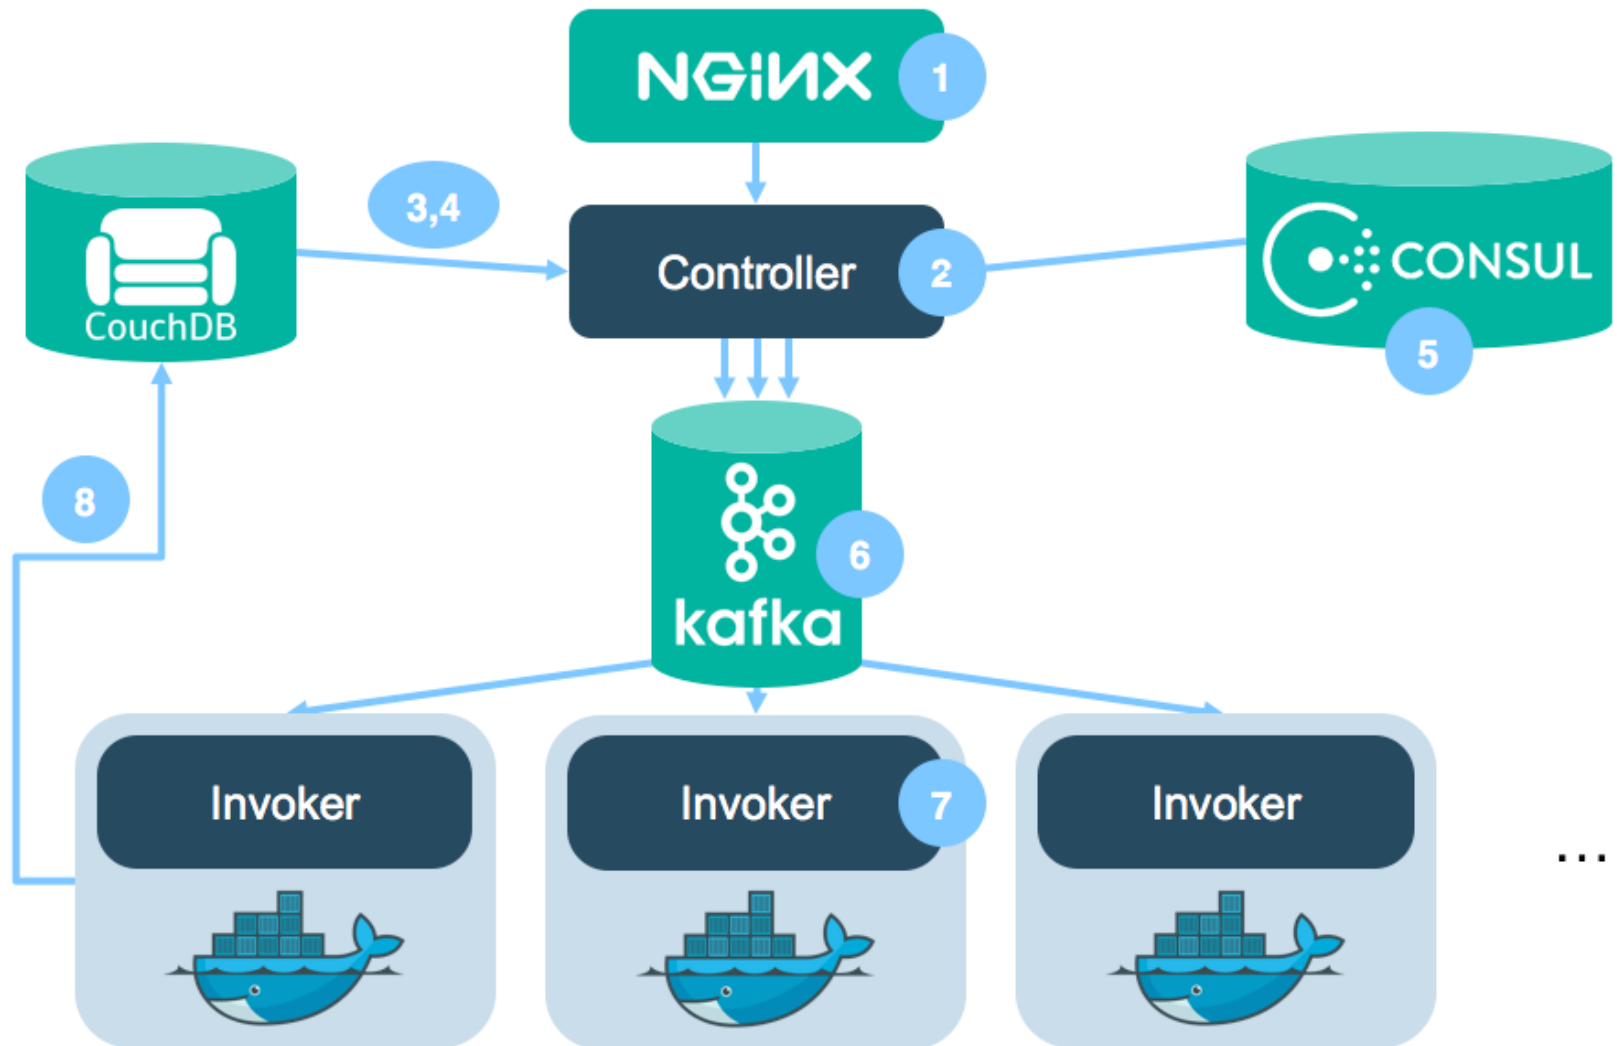
- **Create an action:**
wsk action **create** myfunction Hello-Python.py
- **Invoke an action:**
wsk action **invoke** myfunction --result --param **name** **World**
- **Update an action:**
wsk action **update** myfunction Hello-Python.py

System overview



<https://github.com/apache/incubator-openwhisk/blob/master/docs/about.md>

The internal flow of processing



Advantages of Serverless/FaaS

- Very simple and "cheap" to scale
- Rapid prototyping
- Easy to modify serverless functions
- Pay only for the execution time, not for idle computation time
- Can create applications by composing functions written in different languages

Disadvantages of Serverless

- Harder to avoid vendor lock-in
 - Depend heavily on built in triggers and rules
- Lack of monitoring and debugging tools
- Composition and architecture complexity
- Slow cold-start
- What about stateful computations?
- Harder to predict costs

Next labs

- This week Lab
 - Continue working with Google AppEngine
- Next week Lab
 - Cloud Functions in IBM Bluemix (Managed OpenWhisk service)

Next Lecture

- Overview of Mobile & Cloud Lab research
- Cloud computing challenges
- **NB!** Opportunity to ask additional questions about the exam

References

- Check Amazon videos and webinars at <http://aws.amazon.com/resources/webinars/>
- Mike Roberts, “Serverless Architectures”, <https://martinfowler.com/articles/serverless.html>
- Abel Avram, “FaaS, PaaS, and the Benefits of the Serverless Architecture”, <https://www.infoq.com/news/2016/06/faas-serverless-architecture>
- Apache OpenWhisk - <https://github.com/apache/incubator-openwhisk/blob/master/docs/about.md>
- E. Jonas, J. Schleier-Smith, et. Al. "Cloud programming simplified: A berkeley view on serverless computing." Technical report, University of California, Berkeley, Feb 2019. <https://arxiv.org/abs/1902.03383>